

SF-4: Neutrino Sector Unification from 600-Cell Geometry

Eight Parameters from One Calibration

Conscious Point Physics Flagship Paper Series

Version 0.9 — 9 May 2026 (incorporates ChatGPT review pass 3 feedback on v0.8: JUNO Δm_{21}^2 uncertainty formatting bug fixed (parenthesized correctly as $(7.50 \pm 0.12) \times 10^{-5} \text{ eV}^2$); empirical mass-ratio arithmetic recomputed self-consistently in §3.4 with explicit dual-convention presentation (absolute-mass-comparison 2%/8% and ratio-level 4%/11%); JUNO 2025 bibliography updated from “in preparation” to arXiv:2511.14593; CHANGELOG bookkeeping “Eight specific edits” → “Ten distinct edits”; reviewer convergence on v1.0-promotion-ready)

Thomas Lee Abshier, ND

Claude Opus (Anthropic)

Hyperphysics Institute, Kalispell, Montana

<https://hyperphysics.com>

drthomas007@protonmail.com

OSF DOI: 10.17605/OSF.IO/JXE8D

Abstract

Paper type: Flagship derivation paper. Establishes a parameter-economic derivation of the neutrino sector’s eight observable quantities from the 600-cell substrate machinery developed in the Standard Model Emergence (SM) and Strong Sector (SS) series.

The Standard Model has no internal explanation for neutrino masses, mixing angles, or hierarchy ordering — the eight observable parameters of the neutrino sector are empirical inputs. SF-4 proposes a structurally constrained leading-order derivation of seven of these eight parameters at zero free parameters from a single calibration (m_e , the electron mass) plus the 600-cell substrate geometry and Conscious Point Physics (CPP) primitives. The neutrino mass formula

$$m_{\nu_i} = M_0 \cdot V_{\nu_i}^2 \cdot \sigma_\nu, \quad M_0 = m_e \cdot z/\varphi \approx 3.79 \text{ MeV}, \quad \sigma_\nu = z^{-10} \approx 1.62 \times 10^{-11},$$

uses the same mass quantum M_0 and 600-cell coordination $z = 12$ that anchor the quark and charged-lepton sectors in SM-7 [4], SM-8 [5], SM-9 [6]. The cage-shell vertex counts $V \in \{4, 12, 30\}$ are forced by 600-cell topology (the bonded shells of any 600-cell vertex). The suppression factor $\sigma_\nu = z^{-10}$ is proposed as a structurally constrained leading-order result from substrate walk-dimension primitives for an unbound three-dimensional orbital ZBW mode and matches the empirical absolute mass scale to within 2%. The mass ratio predictions $m_2/m_1 = 9.00$ and $m_3/m_1 = 56.25$ (equivalently, mass-squared ratios $m_2^2/m_1^2 = 81$ and $m_3^2/m_1^2 \approx 3164$) match observation to 4% and 11% on the mass-ratio comparison, with no fitted parameters. The PMNS mixing matrix at zeroth order is exactly the tribimaximal matrix U_{TBM} ,

inherited from the SM-5 K3 spectral derivation [3]; the cage-shell mass mechanism preserves K3-eigenstate alignment by construction once the mass-basis reading is adopted. Normal mass hierarchy is forced by the cage-shell assignment ($\nu_1 \rightarrow V = 4, \nu_2 \rightarrow V = 12, \nu_3 \rightarrow V = 30$). The CP-violating phase δ_{CP} is registered as open per OP-SM-7d and deferred to the electroweak-sector flagship (SF-2).

The paper makes specific zero-parameter quantitative predictions and identifies five clean falsifiers: the JUNO mass-ordering measurement (multi-year program; expected $\sim 3\sigma$ ordering sensitivity within roughly 6 years of full operation) would falsify the cage-shell assignment if inverted hierarchy is confirmed; cosmological tightening to $\Sigma m_\nu < 50$ meV (DESI/CMB-S4) would falsify the $\sigma_\nu = z^{-10}$ prediction (current DESI/Planck bound of $\Sigma m_\nu \leq 72$ meV is one stringent combination, with alternative analyses giving up to ~ 86 meV); the principled direct-mass falsifier — a beta-decay measurement of m_β robustly inconsistent with the predicted $m_\beta \approx 8.7$ meV — is conceptually clean but exceeds near-term experimental sensitivity (Project 8 targets $m_\beta \sim 40$ meV); PMNS deviation from TBM at zeroth order (after SM-5’s higher-order Capotauro corrections) would falsify the K3-eigenmode-identification ansatz that SM-5 introduces and SF-4 inherits; substrate-mechanism deviation from $\sigma_\nu = z^{-2d_{\text{eff}}}$ form would falsify the walk-dimension framework.

The substantive derivations of the suppression mechanism (§4) and the K3-Cage-Shell Consistency Theorem (§5) ship at PARTIAL CLOSURE: structural-physical pictures established with quantitative numerical predictions in hand; theorem-level rigor pending closure of the open theorem-level work catalogued in §10. SF-4 introduces no new fitted parameter beyond the inherited single calibration M_0 , but the cage-shell coupling assignment to specific K3 eigenmodes (Theorem 5.1 clause (iii)) is a new structural-coupling claim whose theorem-level proof remains OPEN-FP-SF-4-2 and is tied to SM-5’s existing open problem on the K3 antibonding-doublet-degeneracy lifting.

Keywords: neutrino mass spectrum, tribimaximal mixing, PMNS matrix, K3 spectral theorem, 600-cell lattice, Conscious Point Physics, walk dimension, Zitterbewegung, cage-shell mass formula, suppression factor, normal hierarchy, JUNO falsifier, OPEN-FP-SF-4-1, OPEN-FP-SF-4-2, Standard Model emergence, parameter-free derivation, flagship paper.

Plain Language Summary: The Standard Model needs eight separate empirical numbers to describe neutrinos: three masses, three mixing angles, a CP-violating phase, and the ordering of the masses. None of these comes out of the SM equations — they are all measured and inserted by hand. This paper proposes a structurally constrained leading-order derivation of seven of those eight numbers from a single measured input (the electron mass) plus the geometry of a four-dimensional shape called the 600-cell. The mass of each neutrino is given by a simple formula: a basic mass scale (≈ 3.79 MeV, computed from the electron mass and the 600-cell’s connectivity number 12), times the square of an integer counting how many vertices a particular geometric shape has, times a small suppression number arising from how the neutrino propagates through the substrate. The integers (4, 12, 30) are forced by the geometry — they are the numbers of vertices in three specific shapes that naturally appear when you measure distances on the 600-cell from any vertex. The suppression number works out to 12^{-10} , which agrees with observation to about 2%. The mixing angles come exactly from a graph theorem about the triangle K_3 , established in an earlier paper (SM-5). The mass ordering — whether the lightest neutrino is closest in mass to the muon neutrino or the tau neutrino — is forced to be the “normal” ordering by the cage-shell assignment, providing a clean test the JUNO experiment will deliver after several years of data (multi-year program targeting roughly 3σ sensitivity within ~ 6 years of full operation). The eighth parameter, the CP phase, is left as an open problem; its derivation belongs in a different paper covering the electroweak sector. Substantive theorem-level rigor for the suppression factor and the cage-shell-

to-K3-eigenmode coupling assignment remains open work catalogued in the closing sections; the paper is explicit about where it ships at theorem level versus partial closure. The paper makes its predictions and falsifiers explicit so that future precision measurements can confirm or kill the framework cleanly.

Contents

1	Introduction	4
1.1	The neutrino sector as a named known-unknown	4
1.2	Strategic posture: strict-C	4
1.3	What this paper delivers	5
1.4	What this paper does not deliver	5
1.5	Position in the SF-line	5
1.6	Claim status ledger	6
2	The SM-5 K3-Eigenmode Foundation	8
2.1	The four-cage taxonomy (SM-1)	8
2.2	The K3 Spectral Theorem (SM-3)	8
2.3	Tribimaximal mixing from K3 eigenmodes (SM-5)	8
2.4	Bound-mode mass-formula machinery (SM-7–SM-9)	9
2.5	Inheritance summary: theorem level versus ansatz level	10
3	The Candidate C Cage-Shell Mass Formula	10
3.1	Statement of the mechanism	10
3.2	Cage-shell assignment in mass basis	10
3.3	First-principles derivation handle for $\alpha = 2$	11
3.4	Mass-ratio predictions at zero free parameters	12
4	The Suppression Mechanism	13
4.1	The walk-dimension framework	13
4.2	Channel enumeration for an unbound 3D orbital ZBW mode	14
4.3	Per-channel suppression: three convergent pictures	15
4.3.1	Picture A: Two-sided DI-bit exchange	15
4.3.2	Picture B: Two ZBW half-cycles per absolute moment	16
4.3.3	Picture C: Edge-straddling coherent state	17
4.3.4	Cross-comparison	17
4.4	Combined result and predictions at zero free parameters	18
4.5	Open theorem-level work (OPEN-FP-SF-4-1)	19
5	The K3-Cage-Shell Consistency Theorem	19
5.1	What the theorem must establish	20
5.2	The mass-basis-vs-flavor-basis clarification	20
5.3	Numerical zeroth-order consistency is exact	20
5.4	The structural-physical question	21
5.5	Route C: Direct distance computation from the lepton position	22
5.5.1	The 600-cell distance-shell structure (verified)	22
5.5.2	$V = 4$ as the tetrahedral subset of shell 1	23

5.5.3	V = 20 exclusion by SM-1 particle-type taxonomy	23
5.6	The K3-eigenmode-to-shell coupling pattern	23
5.7	Route C closure summary and inheritance posture	24
5.8	Open theorem-level work (OPEN-FP-SF-4-2)	25
6	Predictions Summary	26
6.1	Master predictions table	26
6.2	Structural-residual pattern	26
6.3	Comparison to discrete-symmetry A_4 models	27
7	δ_{CP} Posture (Route ii)	27
7.1	Justification of route (ii)	27
7.2	The four candidate handles for route (i)	28
8	Higher-Order Corrections (OP-SM-7d Inheritance)	28
8.1	Inheritance posture	29
8.2	Expected effects of OP-SM-7d closure	29
8.3	What this means for SF-4 v1.0 ship	29
9	Cumulative Falsifier	29
9.1	Direct SF-4 falsifiers	30
9.1.1	Inverted hierarchy (JUNO multi-year program)	30
9.1.2	Principled direct-mass falsifier (longer-timescale)	30
9.1.3	Cosmological tightening to $\Sigma m_\nu < 50$ meV	31
9.2	Framework-level falsifiers	31
9.2.1	PMNS deviation from TBM at zeroth order	31
9.2.2	Substrate-mechanism deviation from $\sigma_\nu = z^{-10}$	32
9.2.3	Partial-failure scenarios and modular falsification	32
9.3	What is not predicted	33
10	Open Theorem-Level Work	33
10.1	Registered open problems	33
10.2	Structural-derivation handle: $\alpha = 2$ closure	34
10.3	Items not addressed in v1.0	34
10.4	Forward roadmap	35
11	Discussion	35
11.1	The programme-level pattern: structural agreement at integer counts as load-bearing signal	35
11.2	Cross-sector implications	36
11.3	Outlook	37

1 Introduction

Paper type: Flagship derivation. Pulls completed sub-derivations from the SM and SS series [1, 2, 3, 4, 5, 6, 7] into an apex synthesis covering the neutrino sector’s eight observable parameters. Strict-C posture: every parameter back to substrate primitives plus a single calibration; the register-as-open card is used judiciously and one or two layers removed from the present problem when the underlying derivation closure lives in another sector.

1.1 The neutrino sector as a named known-unknown

The Standard Model treats the neutrino sector through the same Yukawa-coupling machinery used for charged leptons and quarks, but with masses and mixing angles entered as empirical parameters with no internal explanation. The eight observable parameters of the active-flavor neutrino sector are: three mass eigenvalues $m_{\nu_1}, m_{\nu_2}, m_{\nu_3}$ (or equivalently a lightest-mass scale plus two mass-squared splittings $\Delta m_{21}^2, |\Delta m_{32}^2|$); three mixing angles $\theta_{12}, \theta_{23}, \theta_{13}$ in the PMNS matrix; the CP-violating phase δ_{CP} ; and the mass-squared ordering (normal vs. inverted hierarchy). The Standard Model derives none of these.

In the cage-shell framework of Conscious Point Physics, where charged-lepton and quark masses derive from 600-cell distance shells with single calibration m_e [4, 5, 6], the neutrino sector is the natural next derivation target. The challenge: neutrinos are unbound modes (no central CP anchor; no rigid cage) whereas the existing $V^{7/3}$ mass formula was developed for bound modes (cage-anchored). The mass-derivation extension to unbound modes requires a modified mechanism, identified at Session 39 mechanism selection [9] as Candidate C with exponent $\alpha = 2$.

1.2 Strategic posture: strict-C

This paper adopts the strict-C posture established for the SF-line at Session 37: no compromise on first-principles rigor, every parameter traced back to 600-cell substrate primitives plus single calibration, and the register-as-open card used judiciously and one or two layers removed from the present problem where possible. Where a derivation closure lives in an adjacent sector with its own existing open problems, SF-4 inherits those open problems explicitly rather than re-introducing them as new SF-4 ansatz.

The strict-C posture has consequences for what this paper claims and does not claim:

- **Claims:** the cage-shell mass formula, the cage-shell vertex assignment forced by 600-cell topology, the z^{-10} suppression factor from walk-dimension primitives, the TBM zeroth-order mixing inherited from SM-5, the normal hierarchy from the cage-shell assignment.
- **Inherits as open (existing in adjacent sectors):** SM-5’s K3-eigenmode-identification ansatz [3] (the Open Problem in SM-5 §3 on lifting the K3 antibonding-doublet degeneracy), OP-SM-7d (Capotauro mechanism for higher-order PMNS corrections in the EW sector) for δ_{CP} and for the $\sim 10\%$ deviations from TBM angles.
- **Opens (new SF-4 sub-problems):** OPEN-FP-SF-4-1 (theorem-level Picture A formalization for the suppression mechanism from CPP axioms A1–A11) and OPEN-FP-SF-4-2 (vertex-by-vertex K3-Cage-Shell Consistency Theorem; tied to SM-5’s existing open problem).

1.3 What this paper delivers

The paper delivers, at the structural-derivation level with explicit theorem-level open work:

1. A neutrino mass formula $m_{\nu_i} = M_0 \cdot V_{\nu_i}^2 \cdot \sigma_\nu$ with the same $M_0 = m_e \cdot z/\varphi$ as the quark and charged-lepton sectors, the cage-shell exponent $\alpha = 2$ argued from the bound/unbound boundary (structural argument; theorem-level closure registered as part of OPEN-FP-SF-4-1), and the cage-shell vertex assignment $(V_1, V_2, V_3) = (4, 12, 30)$ forced by 600-cell topology with $V = 20$ excluded by SM-1 particle-type taxonomy.
2. A walk-dimension framework for the suppression factor σ_ν , with $d_{\text{eff}} = 5$ from integer channel enumeration (3 spatial + 1 ZBW phase + 1 orientation) and per-channel suppression z^{-2} from three independent CPP physical pictures, producing $\sigma_\nu = z^{-10} \approx 1.62 \times 10^{-11}$ at 2% empirical match to the absolute neutrino mass scale.
3. A K3-Cage-Shell Consistency Theorem at the SM-5-inheritance level: the mass-basis reading of the cage-shell assignment produces a $\mu\tau$ -symmetric mass operator in flavor basis whose eigenvectors coincide with the TBM directions exactly, and the cage-shell vertex assignment to specific K3 eigenmodes is forced by 600-cell geometry plus the SM-5-ansatzed TBM-direction selection (Route C structural closure).
4. A predictions table covering all eight neutrino parameters, with seven at zero free parameters and the eighth (δ_{CP}) deferred to SF-2 per route (ii). Five clean falsifiers identified.

1.4 What this paper does not deliver

To be honest about scope:

- Theorem-level closure of the suppression mechanism from CPP axioms A1–A11 (Picture A formalization) ships at PARTIAL CLOSURE; rigorous derivation is registered as OPEN-FP-SF-4-1 and is post-v1.0 work.
- Theorem-level closure of OPEN-FP-SF-4-2 at the vertex-by-vertex coupling level is tied to SM-5's existing open problem on lifting the K3 antibonding-doublet degeneracy. SF-4 ships at the SM-5-inheritance level; theorem-level closure requires resolving SM-5's open problem.
- Quantitative δ_{CP} prediction is deferred to SF-2 (route (ii), per OP-SM-7d EW-sector inheritance).
- Majorana versus Dirac character is not specified by the cage-shell mechanism; registered as open.
- Sterile-neutrino predictions are outside the active-flavor scope of this paper; registered as open.

1.5 Position in the SF-line

SF-4 is the active heavy-lift paper in the SF-line, the seven-paper flagship architecture covering all 17 Standard Model particles plus the W^0 CPP-novel prediction [12]. SF-1 (charged leptons), SF-2 (electroweak cage bosons including W^0), SF-3 (quarks) are primarily reframings of established corpus; SF-4 (this paper) is the only SF-line paper requiring substantive new derivation work; SF-5 (strong sector) and SF-6 (electromagnetism: classical/SR/QED unified) are corpus-rich

syntheses; SF-7 (grand unification) synthesizes all six predecessors. The strict-C posture and inheritance discipline established here for SF-4 propagate through the rest of the SF-line.

1.6 Claim status ledger

To prevent overclaiming and to make the paper’s epistemic accounting visible at a single glance, Table 1 summarizes each substantive claim of the SF-4 derivation chain with its current closure status. The ledger explicitly distinguishes (a) results that close at theorem level from CPP primitives or direct topological computation, (b) results inherited as theorem-level statements from prior corpus papers, (c) structural arguments that close numerically but await theorem-level rigor, (d) ansatz-level claims inherited from prior corpus papers, and (e) registered open problems.

Claim	Closure status
$M_0 = m_e \cdot z/\varphi$ as substrate-derived mass scale	THEOREM (inherited from SM-7–SM-9 [4, 5, 6])
$z = 12$ as 600-cell coordination number	THEOREM (inherited from SS-1 [7], SM-1 [1])
Available cage-shell vertex counts $V \in \{4, 12, 30\}$ from any 600-cell reference vertex	THEOREM/COMPUTATION (numerically verified shell sequence §5.5)
$V = 20$ exclusion for the lepton-flavor neutrino sector	INHERITED TAXONOMY (from SM-1 [1] particle-type assignment), <i>not</i> an independent SF-4 theorem
$V^{7/3} \rightarrow V^2$ at the bound/unbound boundary (exponent $\alpha = 2$)	STRUCTURAL ARGUMENT (§3.3); theorem-level closure registered as part of OPEN-FP-SF-4-1
$\sigma_\nu = z^{-10}$ from walk-dimension primitives	PARTIAL CLOSURE (three convergent CPP physical pictures, integer channel enumeration, 2% empirical match; theorem-level rigor v1.0+ via OPEN-FP-SF-4-1 Picture A formalization)
PMNS = U_{TBM} at zeroth order	INHERITED CONDITIONAL THEOREM (SM-5 [3] K3 spectral derivation)
TBM-direction selection in K3 antibonding doublet (which orthonormal pair within the 2D degenerate subspace)	INHERITED ANSATZ (SM-5 explicit open problem on lifting K3 antibonding-doublet degeneracy)
$\nu_1 \leftrightarrow V = 4$, $\nu_2 \leftrightarrow V = 12$, $\nu_3 \leftrightarrow V = 30$ cage-shell-to-K3-eigenmode coupling	PARTIAL CLOSURE (numerical zeroth-order consistency exact via Theorem 5.3; structural closure at SM-5-inheritance level via Route C, §5.5); theorem-level vertex-by-vertex coupling tied to SM-5 open problem; SF-4 introduces no new fitted parameter but registers this as new structural-coupling claim OPEN-FP-SF-4-2
Normal mass hierarchy ordering	FORCED CONSEQUENCE of cage-shell assignment $V_1 < V_2 < V_3$ (§3.2)
Higher-order PMNS corrections (deviations from TBM at $\sim 10\%$)	REGISTER-AS-OPEN (inherited from SM-5 [3] OP-SM-7d Capotauro mechanism; deferred to SF-2 EW-flagship)
δ_{CP} derivation	NOT PREDICTED in v1.0 (route ii deferral to SF-2 per OP-SM-7d, with four candidate handles enumerated for forward reference, §7)
Majorana versus Dirac character	NOT SPECIFIED by the cage-shell mechanism in v1.0; registered as open

Table 1: SF-4 v1.0 claim-status ledger. Each substantive claim of the derivation chain is shown with its closure status (THEOREM, INHERITED, PARTIAL CLOSURE, INHERITED ANSATZ, REGISTER-AS-OPEN, NOT PREDICTED). The ledger makes the framework’s epistemic posture visible at a glance: SF-4 introduces no new fitted parameter beyond the inherited single calibration M_0 , but adds one new structural-coupling claim (the cage-shell-to-K3-eigenmode assignment) at PARTIAL CLOSURE level whose theorem-level closure is OPEN-FP-SF-4-2.

2 The SM-5 K3-Eigenmode Foundation

The neutrino sector derivation in SF-4 builds on five pieces of established CPP corpus: the four-cage taxonomy of SM-1 [1], the K3 Spectral Theorem of SM-3 [2], the tribimaximal mixing result of SM-5 [3], the bound-mode mass-formula machinery of SM-7–SM-9 [4, 5, 6], and the 600-cell substrate foundations of SS-1 [7]. This section recaps the inherited content and is explicit about which inheritance is at theorem level versus ansatz level.

2.1 The four-cage taxonomy (SM-1)

SM-1 establishes the four stable cage geometries available at the 600-cell substrate scale, as bonded shells (or sub-shells) from any 600-cell reference vertex [1]:

- **Tetrahedral cage** ($V = 4$): four vertices in tetrahedral arrangement; subset of shell 1 via the compound-of-5-tetrahedra geometry inscribed in the icosahedron. Hosts the electron and the light quarks.
- **Icosahedral cage** ($V = 12$): full first distance shell at $d^2 = 1/\varphi^2 \approx 0.382$ (in units where 600-cell circumradius is 1). Hosts the charm quark, tau lepton, and Z boson.
- **Dodecahedral cage** ($V = 20$): second distance shell at $d^2 = 1$. Hosts the bottom quark and Higgs boson.
- **Icosidodecahedral shell** ($V = 30$): the $d^2 = 2$ shell, vertex-transitive with degree 4. Hosts the top quark cage.

The four cage assignments propagate into SF-4 via the cage-shell vertex counts that appear in the neutrino mass formula. Direct numerical computation of squared distances from any 600-cell vertex confirms the shell sequence at vertex counts (1, 12, 20, 12, 30, 20, 12, 12, 1) across $d^2 \in [0, 4]$, matching the SM-1 cage assignments at $V = 12, 20, 30$ exactly. Patch 0303 verified this in the working derivation document [11].

2.2 The K3 Spectral Theorem (SM-3)

SM-3 establishes the K3 colour-cage base graph and its spectral theorem [2]. The K3 graph is the equilateral triangle on three colour vertices $\{V_1, V_2, V_3\}$ of the tetrahedral quark cage. Its adjacency matrix

$$A_{K_3} = \begin{pmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix}$$

has eigenvalues $\lambda_+ = +2$ (bonding singlet, multiplicity 1) and $\lambda_- = -1$ (antibonding doublet, multiplicity 2). The bonding eigenvector is $|\phi_+\rangle = (1, 1, 1)^T/\sqrt{3}$. The K3 Spectral Theorem of SM-3 derives the Koide formula $K = (m_e + m_\mu + m_\tau)/(\sqrt{m_e} + \sqrt{m_\mu} + \sqrt{m_\tau})^2 = 2/3$ from the K3 spectral structure, conditional on an open-system thermalisation model whose epistemic status is made explicit via Layer A/B/C decomposition in SM-3 [2].

2.3 Tribimaximal mixing from K3 eigenmodes (SM-5)

SM-5 establishes that the zeroth-order PMNS matrix is exactly the tribimaximal mixing matrix U_{TBM} [3]. The derivation proceeds via two identifications:

- Charged-lepton flavor eigenstates are the K3 vertex states $|V_1\rangle, |V_2\rangle, |V_3\rangle$ (charged leptons couple to the cage through localised vertex occupation).
- Neutrino mass eigenstates are the K3 ZBW Hamiltonian eigenmodes $|\phi_+\rangle, |\phi_-^{(1)}\rangle, |\phi_-^{(2)}\rangle$ (neutrinos, carrying no colour charge, propagate as global oscillation modes of the cage Hamiltonian $\hat{H}_{\text{ZBW}} = \hbar\omega_0 A_{K_3}$).

The PMNS matrix at zeroth order is the change-of-basis between vertex states and K3 eigenvectors:

$$U_{\text{PMNS}}^{(0)} = U_{\text{TBM}} = \begin{pmatrix} \sqrt{2/3} & 1/\sqrt{3} & 0 \\ -1/\sqrt{6} & 1/\sqrt{3} & -1/\sqrt{2} \\ -1/\sqrt{6} & 1/\sqrt{3} & 1/\sqrt{2} \end{pmatrix}, \quad \sin^2 \theta_{12}^{(0)} = \frac{1}{3}, \quad \sin^2 \theta_{23}^{(0)} = \frac{1}{2}, \quad \sin^2 \theta_{13}^{(0)} = 0. \quad (1)$$

The conventional ordering in SM-5, which SF-4 inherits:

$$\nu_1 \leftrightarrow |\phi_-^{(1)}\rangle = \frac{1}{\sqrt{6}}(2, -1, -1)^T, \quad \nu_2 \leftrightarrow |\phi_+\rangle = \frac{1}{\sqrt{3}}(1, 1, 1)^T, \quad \nu_3 \leftrightarrow |\phi_-^{(2)}\rangle = \frac{1}{\sqrt{2}}(0, -1, 1)^T. \quad (2)$$

Remark 2.1 (The SM-5 open problem, inherited by SF-4). *SM-5 §2 explicitly registers the K3-eigenmode-identification of neutrino mass eigenstates as an open problem (“physically motivated but not derived from CPP interaction rules”). A separate aspect of the same open problem: the K3 antibonding doublet $\lambda_- = -1$ has multiplicity 2; any orthonormal pair within this 2D subspace is K3-equivalent. The TBM choice picks specific directions (the $\mu\tau$ -symmetric mode $\phi_-^{(1)}$ and the $\mu\tau$ -antisymmetric mode $\phi_-^{(2)}$), but this selection is not derived from K3 alone. SF-4 inherits both aspects of SM-5’s open problem; SF-4 introduces no new ansatz beyond what SM-5 has already registered.*

2.4 Bound-mode mass-formula machinery (SM-7–SM-9)

SM-7 [4], SM-8 [5], and SM-9 [6] establish the bound-mode mass-formula machinery applied to nuclei and quarks. The mass quantum

$$M_0 = m_e \cdot z/\varphi \approx 3.79 \text{ MeV} \quad (3)$$

where $z = 12$ is the 600-cell coordination number (each vertex has 12 nearest neighbors arranged icosahedrally) and $\varphi = (1 + \sqrt{5})/2$ is the golden ratio, is the substrate-derived mass scale.

Bound-mode masses scale as $M_0 \cdot V^{7/3}$ with V the cage-shell vertex count, decomposing as

$$V^{7/3} = V^2 \cdot V^{1/3} \quad (4)$$

where V^2 is the pair-count component (counting interactions over all pairs of cage vertices) and $V^{1/3}$ is the linear-cage-dimension component (scaling as the linear extent of a 3D cage at fixed local density).

The decomposition (4) is critical for the unbound-mode extension in SF-4: for unbound modes (no rigid cage), the linear-cage-dimension factor $V^{1/3}$ has no operational meaning and drops out, leaving the pair-count component V^2 alone.

Inherited content	Source	Inheritance level
$M_0 = m_e \cdot z/\varphi$ derivation	SM-7–SM-9 [4, 5, 6]	THEOREM
$z = 12$ from 600-cell coordination	SS-1 [7], SM-1	THEOREM
Four-cage taxonomy $V \in \{4, 12, 20, 30\}$	SM-1 [1]	THEOREM
$V^{7/3} = V^2 \cdot V^{1/3}$ decomposition	SM-9 [6]	THEOREM
K3 spectral structure	SM-3 [2]	THEOREM
$U_{\text{PMNS}}^{(0)} = U_{\text{TBM}}$ from K3 eigenvectors	SM-5 [3]	THEOREM (modulo ansatz below)
K3-eigenmode-identification of ν_i mass eigenstates	SM-5 [3] §2	ANSATZ (SM-5 open problem)
TBM-direction selection in K3 antibonding doublet	SM-5 [3] §2	ANSATZ (SM-5 open problem)
Higher-order PMNS corrections (Capotauro mechanism)	OP-SM-7d	REGISTER-AS-OPEN
δ_{CP} derivation	OP-SM-7d (EW sector)	REGISTER-AS-OPEN (route ii)

Table 2: What SF-4 inherits from the SM and SS series corpus, classified by inheritance level. SF-4 introduces no new fitted parameter beyond the single inherited calibration $M_0 = m_e \cdot z/\varphi$; the new structural-coupling claim of cage-shell-to-K3-eigenmode assignment is registered as OPEN-FP-SF-4-2 (§1.6, Table 1).

2.5 Inheritance summary: theorem level versus ansatz level

3 The Candidate C Cage-Shell Mass Formula

3.1 Statement of the mechanism

The Candidate C neutrino mass formula, selected at Session 39 mechanism review [9], is

$$\boxed{m_{\nu_i} = M_0 \cdot V_{\nu_i}^2 \cdot \sigma_\nu} \quad (5)$$

with mass quantum $M_0 \approx 3.79$ MeV given by (3), cage-shell vertex count V_{ν_i} specified in §3.2, exponent $\alpha = 2$ argued from the bound/unbound boundary in §3.3 (structural argument; theorem-level closure part of OPEN-FP-SF-4-1), and suppression factor $\sigma_\nu = z^{-10} \approx 1.62 \times 10^{-11}$ proposed in §4 as a structurally constrained leading-order result.

The formula is structurally continuous with the bound-mode quark/lepton machinery: same M_0 , same 600-cell vertex-count framework, same dimensionless geometry-only structure. The two unbound-regime modifications are: (1) exponent $\alpha = 7/3 \rightarrow 2$ (linear-cage-dimension factor drops out), and (2) introduction of the substrate-derived σ_ν replacing the per-mode μ_q multipliers used in the quark sector [6].

3.2 Cage-shell assignment in mass basis

Proposition 3.1 (Cage-shell assignment, mass basis). *The neutrino mass eigenstates ν_1, ν_2, ν_3 have cage-shell vertex counts*

$$(V_{\nu_1}, V_{\nu_2}, V_{\nu_3}) = (4, 12, 30) \quad (6)$$

corresponding respectively to the tetrahedral cage (subset of 600-cell shell 1), the icosahedral first shell of the 600-cell, and the icosidodecahedral shell at $d^2 = 2$.

Geometric origin of the three V values. The three numbers (4, 12, 30) are not adjusted to fit data; they are the available cage-shell vertex counts forced by 600-cell topology. Specifically: (i) $V = 12$ is the icosahedral first shell — every 600-cell vertex has exactly 12 nearest neighbours arranged at icosahedral angles ($z = 12$ is the 600-cell coordination number); (ii) $V = 4$ is the

tetrahedral inscribed sub-cage of shell 1 — the icosahedron contains five inscribed regular tetrahedra (the famous compound-of-five-tetrahedra), each using 4 of the 12 first-shell vertices, and the lepton-flavor structure picks out one such tetrahedron via SM-1’s four-cage taxonomy; (iii) $V = 30$ is the icosidodecahedral shell at squared distance $d^2 = 2$ — the third bonded shell from any 600-cell reference vertex, with 30 vertices arranged in 15 antipodal pairs. Detailed verification of the shell sequence is given in §5.5 (Table 5); the SM-1 particle-type taxonomy of which cage hosts which Standard Model particle (and consequently which shells the neutrino sector can couple to) is established in [1].

Mass basis vs flavor basis (key clarification). The cage-shell assignment (6) is in *mass basis* (the K3-eigenmode basis from SM-5 [3], equation (2)), *not* flavor basis (the K3 vertex basis $|V_\alpha\rangle$ corresponding to charged-lepton flavor labels). The mass-basis reading is forced by the requirement that $U_{\text{PMNS}}^{(0)} = U_{\text{TBM}}$ continue to hold; the alternative flavor-basis reading would force PMNS = identity, contradicting both SM-5 and observation. This is the load-bearing observation of the K3-Cage-Shell Consistency Theorem (§5, in particular §5.2); it is highlighted here at first appearance of (6) because the entire derivation chain of SF-4 hangs on the mass-basis reading.

Why $V = 20$ is excluded. The four-cage taxonomy [1] assigns $V = 20$ (the dodecahedral second shell) to the bottom quark and Higgs boson. The lepton-flavor neutrino sector does not couple to this shell; SF-4’s cage-shell assignment uses the three available shells $\{4, 12, 30\}$ from the lepton-cage substrate. The $V = 20$ exclusion is forced by the SM-1 particle-type taxonomy, not by SF-4 fitting.

3.3 First-principles derivation handle for $\alpha = 2$

The bound-mode formula has exponent $V^{7/3}$ derived in SM-9 via the decomposition (4): $V^{7/3} = V^2 \cdot V^{1/3}$, where V^2 counts interactions over all cage-vertex pairs (the pair-count component) and $V^{1/3}$ is the linear-cage-dimension component capturing the spatial extent of a rigid 3D cage at fixed local density.

For unbound modes (no central CP anchor; no rigid cage by hypothesis), no rigid cage exists; the linear-cage-dimension factor $V^{1/3}$ has no operational meaning. The exponent reduces to V^2 :

$$V^{7/3} \rightarrow V^2 \quad (\text{unbound regime}). \quad (7)$$

This is not a fit. The exponent $\alpha = 2$ is forced by the structural difference between bound and unbound modes — exactly the boundary across which $V^{1/3}$ does or does not contribute to the mass formula.

Operator-level picture. In CPP terms, the neutrino mass operator is proportional to the substrate pair-interaction count summed over the cage-shell vertices. For the bound-mode quark/lepton sectors this operator is multiplied by a linear-cage-dimension factor capturing the spatial extent of a rigid 3D cage at fixed local density. For the unbound neutrino mode, the operator retains its V^2 pair-count scaling, while the geometric extent term has no CPP-operational definition (no rigid cage exists for an unbound 3D orbital ZBW configuration; there is no spatial extent to characterize). The structural argument for $\alpha = 2$ is that the operator decomposition $M_0 \cdot V^{7/3} = M_0 \cdot V^2 \cdot V^{1/3}$ has its second factor $V^{1/3}$ undefined for unbound modes, so the leading non-zero contribution is the pair-count scaling alone.

Theorem-level closure of (7) from CPP axioms A1–A11 is registered as part of OPEN-FP-SF-4-1 (suppression mechanism Picture A formalization) and is post-v1.0 work.

3.4 Mass-ratio predictions at zero free parameters

With $\alpha = 2$ and the assignment (6), the mass ratios are

$$\frac{m_{\nu_2}}{m_{\nu_1}} = \left(\frac{12}{4}\right)^2 = 9.00, \quad \frac{m_{\nu_3}}{m_{\nu_1}} = \left(\frac{30}{4}\right)^2 = 56.25. \quad (8)$$

The observed mass-squared splittings $\Delta m_{21}^2 = (7.50 \pm 0.12) \times 10^{-5} \text{ eV}^2$ (JUNO 2025 first physics results [18]; consistent with NuFIT 6.0 [13] value $7.49 \times 10^{-5} \text{ eV}^2$) and $|\Delta m_{32}^2| = 2.513 \times 10^{-3} \text{ eV}^2$ (NuFIT 6.0) determine empirical absolute mass values once the lightest mass m_1 is fixed.

Adopting the lightest-massless approximation $m_1 \rightarrow 0$ as the comparison convention (accurate to better than 1% since $m_1 \ll \sqrt{\Delta m_{21}^2}$ on any reasonable assumption about the absolute scale), the empirical absolute masses are

$$m_2 \rightarrow \sqrt{\Delta m_{21}^2} \approx 8.66 \text{ meV}, \quad m_3 \rightarrow \sqrt{|\Delta m_{31}^2|} \approx 50.9 \text{ meV}, \quad (9)$$

where $|\Delta m_{31}^2| = \Delta m_{21}^2 + |\Delta m_{32}^2| \approx 2.588 \times 10^{-3} \text{ eV}^2$. *Comparison conventions for the SF-4 ratio predictions.* The SF-4 mass-ratio predictions (8) can be compared to observation under two equivalent conventions:

1. **Absolute-mass comparison** (Table 4, §4.4): combine the cage-shell ratio predictions with the SF-4 absolute-scale prediction $m_1 = M_0 V_1^2 \sigma_\nu \approx 0.98 \text{ meV}$ (§4) to give SF-4 absolute predictions $m_2 = 8.81 \text{ meV}$ and $m_3 = 55.1 \text{ meV}$. Comparison to the empirical (9) gives 2% match on m_2 and 8% match on m_3 .
2. **Ratio-level comparison** (under implicit reference normalization $m_1 = 1 \text{ meV}$, close to but not identical to the SF-4 prediction $m_1 \approx 0.98 \text{ meV}$): treating the empirical mass values 8.66 and 50.9 meV from (9) as effective ratios under that reference, the SF-4 ratio predictions 9.00 and 56.25 give 4% and 11% match levels. Equivalently, adopting the SF-4 $m_1 = 0.98 \text{ meV}$ directly: empirical ratios become $m_2/m_1 \approx 8.9$ and $m_3/m_1 \approx 52$, giving $\sim 1\%$ and $\sim 8\%$ match levels.

The two conventions show the same underlying residual pattern: tight on m_2/ν_2 ($\sim 1\text{--}4\%$), modest on m_3/ν_3 ($\sim 8\text{--}11\%$). The 2% and 8% absolute-mass match levels are the convention used in Table 4 and §6.2 below; the 4% and 11% ratio-level match levels appear in summary statements where the dimensionless ratio is the primary quantity. Both are valid views on the same structural residual.

The structural residuals (2–11% across the two conventions) reflect the leading-order V^2 approximation and are expected to be *structural* rather than *statistical*: the same V^2 scaling assumption produces them across both ratios simultaneously, and the OP-SM-7d Capotauro mechanism (§8) is the natural lifting mechanism. They recur as the same residual pattern in the absolute mass predictions (§4.4).

The comparison in (8) is at the *mass ratio* level, not the mass-squared splitting level. The cage-shell formula (5) predicts mass eigenvalues directly via V^2 scaling; the corresponding mass-squared ratios are $m_2^2/m_1^2 = 81$ and $m_3^2/m_1^2 \approx 3164$, and the mass-squared splitting ratio $\Delta m_{31}^2/\Delta m_{21}^2 = (m_3^2 - m_1^2)/(m_2^2 - m_1^2) \approx 39.5$ (predicted) versus ~ 34.1 (empirical from NuFIT

6.0) at the 13% match level. The mass-ratio comparison in (8) is the cleanest direct test of the V^2 scaling structure; the mass-squared splitting ratio is the historically standard observable but reflects the same underlying structural prediction.

Two distinct structural claims. The neutrino sector predictions partition into two structurally distinct classes: (a) the *mass ratios* $m_2/m_1 = 9$ and $m_3/m_1 = 56.25$ are zero-parameter predictions of the V^2 scaling forced by 600-cell topology alone; (b) the *absolute mass scale* (and consequently $m_{\nu_1} \approx 0.98$ meV, $\Sigma m_\nu \approx 64.9$ meV) is a separate structural claim via the suppression factor $\sigma_\nu = z^{-10}$, which carries its own PARTIAL CLOSURE status as the substantive content of §4. The mass-ratio claim is robust to closure level of σ_ν ; the absolute-scale claim is conditional on the OPEN-FP-SF-4-1 closure path. Master predictions Table 6 (§6) and the claim-status ledger (Table 1, §1.6) keep this distinction visible.

The structural agreement at zero free parameters is the load-bearing signal; the residuals are the framework’s intrinsic precision at the leading-order V^2 approximation. Higher-order corrections via OP-SM-7d would tighten these.

4 The Suppression Mechanism

The cage-shell mass formula (5) is anchored at $M_0 \approx 3.79$ MeV, which is twelve orders of magnitude above the empirical neutrino mass scale of $\sim 1\text{--}10^2$ meV. The suppression factor σ_ν closes that gap. This section proposes a structurally constrained leading-order derivation of $\sigma_\nu = z^{-10} \approx 1.62 \times 10^{-11}$ from substrate walk-dimension primitives via integer channel enumeration ($d_{\text{eff}} = 5$) and per-channel suppression $1/z^2$ established through three independent and convergent CPP physical pictures. Theorem-level rigor from CPP axioms A1–A11 is registered as OPEN-FP-SF-4-1 and is post-v1.0 work.

The headline result, stated in advance:

$$\boxed{\sigma_\nu = z^{-2d_{\text{eff}}} = 12^{-10} \approx 1.62 \times 10^{-11}} \quad (10)$$

matching the empirical absolute-mass-scale target of $\approx 1.59 \times 10^{-11}$ to within 2%.

4.1 The walk-dimension framework

In CPP, an unbound ZBW configuration propagating through the substrate must coherently maintain its mode characteristics across each absolute moment. Each moment, the mode interacts with the surrounding 600-cell substrate via DI-bit exchanges; each interaction samples a finite set of substrate options, with the mode’s next state stochastically distributed across those options.

Definition 4.1 (Walk channel). *A walk channel is an independent substrate-information degree of freedom whose state must be coherently maintained per absolute moment. Each channel contributes a multiplicative dilution factor $1/N_{\text{channel}}$ to the mode’s coherence per moment, where N_{channel} is the substrate-information count for that channel.*

Definition 4.2 (Walk dimension). *The walk dimension d_{eff} is the count of independent walk channels for the mode. Under the simplifying assumption that all channels share a common per-channel suppression strength σ_{channel} , the total per-moment suppression is*

$$\sigma_\nu = \sigma_{\text{channel}}^{d_{\text{eff}}} \quad (11)$$

The walk-dimension framework formalizes the substrate-coupling structure that produces large suppression of unbound-mode propagation. The framework is structurally agnostic about which substrate primitive supplies σ_{channel} per channel and which channels are independent; both questions are answered below for the unbound 3D orbital ZBW configuration relevant to neutrinos.

Bound versus unbound boundary. For *bound modes* (charged leptons, quarks), the mode sits in a cage of CPs that pins its state. The cage-resonance condition locks the ZBW phase, the cage symmetry locks orbital orientation, and the central CP locks spatial position. None of the channels samples the substrate freely — all are substrate-locked. Therefore $d_{\text{eff}} = 0$ and $\sigma_{\nu} = 1$ for bound modes, recovering the bound-mode mass formula $m = M_0 \cdot V^{7/3}$ of SM-9 [6] without additional suppression.

For *unbound modes* (neutrinos), no central CP anchor exists and no rigid cage forms. Each channel samples the substrate freely per absolute moment. The walk dimension d_{eff} is the count of these free channels. §4.2 enumerates them; §4.3 proposes the per-channel suppression strength via three convergent physical pictures.

4.2 Channel enumeration for an unbound 3D orbital ZBW mode

The neutrino is identified [8] as an unbound 3D orbital ZBW configuration of dipole-pair structures with no central CP anchor. The walk channels of such a mode are enumerated as follows.

Spatial position channels (3). An unbound 3D orbital ZBW mode propagates through 3-dimensional space. Its spatial position is unconstrained by any cage. Each spatial axis (x, y, z) is an independent channel: per absolute moment, the mode’s position along that axis is stochastically distributed across the substrate-information set for that direction. Channel count: 3. This count is essentially definitional — “3D orbital” means three spatial dimensions of free position, by hypothesis.

ZBW oscillation phase channel (1). The ZBW oscillation has an angular phase θ that advances per absolute moment. For a bound mode, the phase is locked to the cage-resonance condition: the cage CPs supply the boundary structure that fixes θ at substrate-coherent values, and the mode’s phase is not free to sample the substrate. For an unbound mode, no such boundary exists: the phase is sampled from the substrate’s phase-information content per moment. Channel count: 1.

Remark 4.3 (Coarse-grained vs. stochastic phase advance). *In the standard quantum-mechanical free-particle picture, the phase advances deterministically at $\omega = mc^2/\hbar$ — there is no stochastic sampling. In the CPP picture, the deterministic-looking advance is the coarse-grained outcome of substrate-internal DI-bit exchanges that do sample the substrate per moment; the phase is not in the strict sense “decided” per moment but is “transmitted” through substrate channels that contribute to coherence dilution. This is the relevant sense in which it counts as a walk channel. A rigorous derivation closing this requires CPP-internal accounting of phase-information transmission per absolute moment; see §10 for OPEN-FP-SF-4-1 work.*

Orbital orientation channel (1). The ZBW orbital plane (equivalently, the angular-momentum direction) is the geometric orientation of the unbound mode’s circulation. For a bound mode, orientation is locked by the cage symmetry; for an unbound mode, the orbital plane orients freely and samples substrate orientation states per moment. Channel count: 1.

In CPP, fermion spin arises from the inner ZBW orbital at twice the outer-orbital frequency (per the existing CPP postulate set). Spin and orbital orientation are physically the same channel — the inner-outer 2:1 frequency relationship locks them to a single geometric direction. Counting them as one channel rather than two is the right discipline.

Total walk dimension.

$$d_{\text{eff}} = \underbrace{3}_{\text{spatial}} + \underbrace{1}_{\text{ZBW phase}} + \underbrace{1}_{\text{orientation}} = 5. \quad (12)$$

This is a leading-order count from integer channel enumeration. Sub-leading contributions may arise from partial-binding effects (the K3-eigenstructure constraint partially aligns one channel with substrate eigenmodes; cf. §5), finer channel decomposition if orbital orientation and intrinsic spin partially decouple in the unbound regime, or substrate-count-base refinements. These are queued as v1.0+ work in OPEN-FP-SF-4-1.

4.3 Per-channel suppression: three convergent pictures

The empirical 2% match with $\sigma_\nu = z^{-2d_{\text{eff}}}$, $z = 12$, $d_{\text{eff}} = 5$ requires $\sigma_{\text{channel}} = 1/z^2$ per walk channel. Three candidate physical pictures all give the same numerical answer; they differ in which CPP primitive does the work, and in which closure path is most natural for theorem-level rigor.

The discipline at this stage is to lay them out as candidates rather than pick one prematurely. The underlying physics has too many degrees of freedom and not enough independent constraints to single-pick at the v1.0 level. Each picture is internally consistent and consistent with CPP's existing primitives; the eventual selection will come from cross-checking against (a) other unbound-mode physics where one picture predicts differently from another, and (b) which closure path actually goes through under axioms A1–A11.

4.3.1 Picture A: Two-sided DI-bit exchange

Story. Per absolute moment, the substrate's fundamental information-transmission unit is the DI-bit exchange. A DI-bit exchange has a send-side and a receive-side: a CP at vertex v_i releases information into a substrate channel toward one of its $z = 12$ neighbors, and the receiving CP at the destination vertex accepts information from one of its $z = 12$ neighbors. For a walk channel of an unbound mode to maintain coherence across the moment, both the send-direction and the receive-direction must align with the channel's required orientation simultaneously.

Counting. Per channel per moment:

- Send-side has z free options (one per neighbor of the source vertex).
- Receive-side has z free options (one per neighbor of the destination vertex).
- The channel is coherent only when both sides hit their required state.

The send and receive choices are independent per channel per moment in the leading-order picture. The probability of coherent transmission is therefore

$$\sigma_{\text{channel}}^{(A)} = \frac{1}{z} \cdot \frac{1}{z} = \frac{1}{z^2}. \quad (13)$$

Why bound modes don't carry this factor. For a bound mode in a cage, the cage CPs supply specific boundary conditions on both source and receive sides — both are pinned by the cage geometry to specific values. There is no free choice on either side, so the $(1/z) \cdot (1/z)$ factor reduces to $(1/1) \cdot (1/1) = 1$ and $\sigma_\nu = 1$ for bound modes, as required.

Closure path from A1–A11. This picture anchors most directly on DI-bit exchange as a substrate primitive. Theorem-level closure requires (a) formally defining “send-side” and “receive-side” of a DI-bit exchange in axiomatic CPP, (b) proving that for an unbound mode each side samples freely from z options, and (c) proving that channel coherence is the AND of both sides hitting the required state. Each is a tractable step in the existing axiom system. This is the most CPP-axiomatic picture and the cleanest theorem-level closure path.

Why it might not be true. The argument requires that the send and receive choices are *independent* per channel per moment. If the substrate's correlation structure ties them — e.g., if the destination CP's receive-side is determined by the source CP's send-side via some substrate-level coherence mechanism that isn't itself a free choice — then the effective per-channel suppression collapses to $1/z$ and the framework needs reconfiguration. This is the most important sub-question for theorem-level closure of OPEN-FP-SF-4-1.

4.3.2 Picture B: Two ZBW half-cycles per absolute moment

Story. The fermion ZBW structure in CPP has an inner orbital at twice the outer-orbital frequency. Each absolute moment contains one full inner-orbital cycle and a half-cycle of the outer orbital; equivalently, two inner-orbital half-cycles per moment. Each half-cycle independently samples one substrate direction (one of z neighbors). Channel coherence requires both half-cycle choices to align with the channel's required direction.

Counting. Per channel per moment:

$$\sigma_{\text{channel}}^{(B)} = \frac{1}{z} \cdot \frac{1}{z} = \frac{1}{z^2}. \quad (14)$$

Why bound modes don't carry this factor. For bound modes the inner-outer 2:1 frequency ratio is locked to the cage resonance — both half-cycles are pinned to specific phase values by the cage geometry. No free sampling on either half-cycle, so $\sigma_\nu = 1$.

Closure path from A1–A11. Picture B anchors on the existing 2:1 frequency convention for fermion structure. Theorem-level closure requires showing that for unbound modes, the two half-cycles are independently sampling — i.e., that the unbound regime preserves the 2:1 structure but releases the cage-imposed coupling between half-cycles. Both pieces look tractable but require explicit work.

Why it might be true and why it might fail. The 2:1 frequency convention is already a CPP postulate (per the existing fermion-spin machinery). It would be elegant if the same structural feature that produces fermion spin (inner-outer orbital coupling at 2:1) also produces the per-channel z^{-2} suppression for unbound modes — a unified explanation across spin and mass. The picture depends specifically on the 2:1 convention being the right number for the unbound regime. If the underlying CPP physics doesn't actually produce a sharp

two-half-cycles-per-moment structure in the unbound regime — e.g., if the half-cycle counting is an artifact of the bound-mode cage geometry that doesn't survive into the unbound regime — the picture fails. Picture A doesn't have this dependency.

4.3.3 Picture C: Edge-straddling coherent state

Story. An unbound mode's coherent state is not localized at a single 600-cell vertex but straddles an edge of the polytope — the mode's wavefunction has support on a pair of adjacent vertices simultaneously. Per absolute moment, the mode transitions from one edge to a neighboring edge; both endpoints of the new edge must be in the right relationship to the channel's direction.

Counting. Per channel per moment:

$$\sigma_{\text{channel}}^{(C)} \approx \frac{1}{z^2}. \quad (15)$$

Source endpoint of the new edge has z options; destination endpoint has z options constrained to form a valid 600-cell edge. Approximately z^2 pair configurations are available; channel coherence requires the specific edge aligning with the channel direction.

Why bound modes don't carry this factor. Bound modes are vertex-localized (anchored at the central CP), not edge-straddling. Picture C's mechanism doesn't apply to bound modes; the bound mode follows a different scaling that produces $\sigma_\nu = 1$ at the cage-pinned configuration.

Closure path from A1–A11. Picture C requires a genuinely new postulate or a derivation: that unbound modes are edge-straddling rather than vertex-localized. This is not currently in the CPP postulate set. Closure would require axiomatic addition or a derivation from existing primitives (perhaps energy-minimization arguments showing edge-straddling is preferred for unbound modes). Less direct than Pictures A and B.

Why it might be true and why it might fail. Edge-straddling has natural connections to gauge-field structure in CPP — gauge bosons ($W^\pm$, W^0 , Z , gluons) are sometimes pictured as living on edges of the polytope rather than vertices, anticipating SF-2 work. If unbound fermions and gauge bosons share an edge-straddling structure, there is a unifying picture that helps the EW-sector flagship later. However, no current CPP postulate or derivation supports edge-straddling for fermions specifically. Introducing it as a new postulate to explain the neutrino mass scale is exactly the kind of move strict-C wants to avoid; Picture C should be adopted only if Pictures A and B both fail closure.

4.3.4 Cross-comparison

The three pictures decohere when applied to other unbound-mode physics in CPP. For instance, Picture B's prediction depends specifically on the 2:1 inner-outer fermion-spin convention and would not apply (without modification) to scalar unbound modes; Picture A applies generically to any unbound DI-bit-mediated propagation. Cross-pollination with adjacent CPP physics (other unbound-mode mass scales, free-particle propagators, light propagation) is one route to discriminate among the three pictures empirically. This work is queued as v1.0+ refinement.

	Picture A	Picture B	Picture C
Anchors on	DI-bit exchange (substrate primitive)	2:1 frequency convention (fermion structure)	Edge-straddling (new postulate / derivation)
Numerical result	$\sigma_{\text{channel}} = 1/z^2$	$\sigma_{\text{channel}} = 1/z^2$	$\sigma_{\text{channel}} \approx 1/z^2$
Closure path under A1–A11	Most direct (DI-bit already a CPP primitive)	Direct (relies on 2:1 convention)	Requires new postulate or derivation
Programme-level coherence	Standalone (specific to mass)	Connects to spin (unified explanation)	Connects to gauge sector (SF-2 alignment)
v1.0 status	LEADING candidate for theorem-level closure	Live alternative	Speculative; adopt if A and B both fail

Table 3: Cross-comparison of the three candidate physical pictures for $\sigma_{\text{channel}} = 1/z^2$. All three converge on the same numerical answer and all three correctly predict $\sigma_\nu = 1$ for bound modes. The robustness of the numerical result across three independent pictures is itself a positive structural signal: $\sigma_\nu = z^{-2d_{\text{eff}}}$ is not contingent on a specific mechanism choice. Picture A is the priority closure path for theorem-level rigor at v1.0+; Pictures B and C remain available as alternatives if Picture A surfaces an obstruction.

4.4 Combined result and predictions at zero free parameters

With Picture A (or B, or C) supplying $\sigma_{\text{channel}} = 1/z^2$ per walk channel, and the channel enumeration of §4.2 giving $d_{\text{eff}} = 5$:

$$\sigma_\nu = \prod_{\text{channels}} \sigma_{\text{channel}} = \left(\frac{1}{z^2}\right)^{d_{\text{eff}}} = z^{-2d_{\text{eff}}} = 12^{-10} \approx 1.62 \times 10^{-11}, \quad (16)$$

matching the boxed headline result (10). Substituted into the cage-shell mass formula (5) with the cage-shell assignment (6):

Quantity	SF-4 prediction	Empirical / observational	Match
σ_ν	1.62×10^{-11}	$\sim 1.59 \times 10^{-11}$ (target)	within 2%
m_{ν_1} (lightest, NH)	0.98 meV	≤ 5 meV (cosmological + splittings)	within consistency window
m_{ν_2}	8.81 meV	~ 8.66 meV (from Δm_{21}^2)	within 2%
m_{ν_3}	55.1 meV	~ 50.9 meV (from Δm_{32}^2)	within 8%
Σm_ν	64.9 meV	≤ 72 meV (DESI/Planck combined)	within bound

Table 4: Neutrino masses predicted by the cage-shell formula (5) with $M_0 = 3.79$ MeV, $V \in \{4, 12, 30\}$, and $\sigma_\nu = 1.62 \times 10^{-11}$. All predictions at zero free parameters. The 2% match in σ_ν , the 2% match in m_{ν_2} , and the 8% match in m_{ν_3} are structural agreement at the leading-order V^2 approximation; the $\sim 8\%$ residual in m_{ν_3} tracks the same residual pattern as the V^2 ratio prediction (8).

Note on the $z = 12$ factor. The integer 12 entering the suppression formula is the 600-cell coordination number — each vertex has 12 nearest neighbors arranged icosahedrally. This is the same 12 that enters the mass quantum $M_0 = m_e \cdot z/\varphi$ in (3) via SM-7–SM-9 [4, 5, 6], and the same 12 that is the vertex count of the icosahedral cage shell $V_{\nu_2} = 12$ in the cage-shell assignment (6). The 12 is not a fitted parameter; it is a 600-cell topological invariant that recurs across multiple structural roles in the SF-4 derivation.

Cosmological-bound sanity check. The total absolute scale prediction

$$\Sigma m_\nu = m_{\nu_1} + m_{\nu_2} + m_{\nu_3} = M_0 \cdot \sigma_\nu \cdot (V_1^2 + V_2^2 + V_3^2) = M_0 \cdot \sigma_\nu \cdot 1060 \approx 64.9 \text{ meV} \quad (17)$$

is computed at zero free parameters from the inherited $M_0 \approx 3.79 \text{ MeV}$, the topologically forced cage shells $V \in \{4, 12, 30\}$ (sum-of-squares $16 + 144 + 900 = 1060$), and the structurally constrained $\sigma_\nu = z^{-10} \approx 1.62 \times 10^{-11}$. The result $\Sigma m_\nu \approx 64.9 \text{ meV}$ is below the most stringent current cosmological combination ($\Sigma m_\nu \leq 72 \text{ meV}$ from DESI BAO + Planck 2018 [15, 16]) with limited margin, and well below the alternative-analysis bound ($\sim 86 \text{ meV}$ from Planck PR4 + supernova datasets [14]). **Note that no empirical input was used in deriving (17);** the absolute-scale prediction is purely structural via σ_ν and the empirical-comparison column of Table 4 is a check of the framework against observation rather than a tuning of the framework to observation. The cosmological-bound sanity check thus also provides one of the cleanest near-term framework falsifiers (§9.1.3).

4.5 Open theorem-level work (OPEN-FP-SF-4-1)

The structural-physical picture established in §§4.1–4.3 is at PARTIAL CLOSURE: three independent CPP physical pictures converge on $\sigma_{\text{channel}} = 1/z^2$, integer channel enumeration is established, the combined $\sigma_\nu = z^{-10}$ matches empirical at 2%, and the bound/unbound boundary is consistent across all three pictures. Theorem-level closure from CPP axioms A1–A11 is registered as OPEN-FP-SF-4-1, with the following sub-goals:

1. Picture A formalization (priority): rigorously establish (a) that DI-bit send-side and receive-side are independent in the unbound regime, (b) that channel coherence is the AND of both sides aligning, (c) the per-moment $1/z^2$ result from these two facts.
2. Independence verification: the leading risk for Picture A is that send and receive choices might be substrate-correlated; rigorous derivation must rule this out or quantify the correction.
3. Channel-count rigor: the integer count $d_{\text{eff}} = 5$ is leading-order; sub-leading corrections from partial-binding effects (K3-eigenstructure constraint partially aligns one channel with substrate eigenmodes; cf. §5) must be quantified.
4. Bound/unbound boundary derivation: $V^{7/3} \rightarrow V^2$ in (7) is currently structurally argued; rigorous derivation from A1–A11 closes the $\alpha = 2$ exponent simultaneously.

Theorem-level closure is post-v1.0 work. The SF-4 paper ships at PARTIAL-CLOSURE level for OPEN-FP-SF-4-1, with explicit pointers to this open work.

5 The K3-Cage-Shell Consistency Theorem

The cage-shell mass formula (5) attaches three numerical mass values $\{16, 144, 900\} \cdot M_0 \sigma_\nu$ to the three neutrino species via the cage-shell vertex assignment $V \in \{4, 12, 30\}$. SM-5 [3] attaches three TBM mixing-amplitude eigenstates to the same three neutrino species via the K3 ZBW Hamiltonian eigenmode identification. Both must hold simultaneously: the SF-4 mass formula must produce mass eigenstates that align with SM-5’s K3-eigenmode identification, otherwise the SM-5 PMNS theorem is silently broken.

The K3-Cage-Shell Consistency Theorem establishes this alignment. The numerical part is exact at zeroth order, by construction once the mass-basis reading of (6) is adopted (§5.2); the

structural-physical part closes at SM-5-inheritance level via direct 600-cell distance-shell computation (§5.5). SF-4 introduces no new fitted parameter beyond the inherited single calibration M_0 , but the assignment of specific cage-shell V values to specific K3 eigenmodes (Theorem 5.1 clause (iii)) is itself a new structural-coupling claim whose theorem-level proof is registered as OPEN-FP-SF-4-2 and is tied to SM-5's existing open problem on the K3 antibonding-doublet-degeneracy lifting.

5.1 What the theorem must establish

Theorem 5.1 (K3-Cage-Shell Consistency, structural-numerical level). *Let $\hat{M} : \mathcal{H}_{K_3} \rightarrow \mathcal{H}_{K_3}$ be the neutrino mass operator on the three-dimensional K3 Hilbert space, defined by the cage-shell mass formula (5) with assignment (6). Then:*

- (i) *The eigenstates of $\hat{M}$ coincide with the K3 ZBW Hamiltonian eigenmodes $|\phi_-^{(1)}\rangle, |\phi_+\rangle, |\phi_-^{(2)}\rangle$ defined in (2).*
- (ii) *The PMNS mixing matrix in (1) is recovered exactly: $U_{\text{PMNS}}^{(0)} = U_{\text{TBM}}$, with $\sin^2 \theta_{12} = 1/3$, $\sin^2 \theta_{23} = 1/2$, $\sin^2 \theta_{13} = 0$.*
- (iii) *The cage-shell V assignment $(V_{\nu_1}, V_{\nu_2}, V_{\nu_3}) = (4, 12, 30)$ to specific K3 eigenmodes is forced by 600-cell topology and SM-1 particle-type taxonomy, modulo SM-5's existing open problem on lifting the K3 antibonding-doublet degeneracy.*

Clauses (i) and (ii) close exactly at zeroth order, by direct calculation in the K3-eigenmode basis with the mass-basis reading of (6) adopted (§§5.2, 5.3). Clause (iii) closes structurally at the SM-5-inheritance level via Route C (§5.5), with theorem-level rigor on the antibonding-doublet split tied to SM-5's existing open problem.

5.2 The mass-basis-vs-flavor-basis clarification

The cage-shell assignment $(V_{\nu_1}, V_{\nu_2}, V_{\nu_3}) = (4, 12, 30)$ is in *mass basis* — V values attach to neutrino mass eigenstates (the K3 ZBW eigenmodes of SM-5, equation (2)), not to flavor states (the K3 vertex states $|V_\alpha\rangle$ corresponding to charged-lepton flavor labels).

This clarification, the load-bearing observation of Session 42 work [11] §3, is forced by the requirement that $U_{\text{PMNS}}^{(0)} = U_{\text{TBM}}$ continue to hold. Under the alternative flavor-basis reading, the mass operator in flavor basis would be diagonal with eigenvalues (16, 144, 900) on the flavor states ($|V_1\rangle, |V_2\rangle, |V_3\rangle$), mass eigenstates would coincide with flavor states, and the PMNS matrix would be the identity — contradicting both SM-5 and observation entirely. Only the mass-basis reading is consistent.

The mass-basis reading is also physically natural: the SF-4 mass formula $m_{\nu_i} = M_0 \cdot V_{\nu_i}^2 \cdot \sigma_\nu$ specifies mass eigenvalues, and mass eigenvalues are properties of mass eigenstates. The cage-shell V is a property of how the mass eigenstate propagates through the substrate, not a property of the flavor-state-to-mass-state mixing matrix.

5.3 Numerical zeroth-order consistency is exact

With the mass-basis reading clarified, Theorem 5.1 clauses (i) and (ii) close by direct calculation in the K3-eigenmode basis. Construct the V^2 operator on the K3-eigenmode basis with

eigenvalues $V^2 = (16, 144, 900)$ assigned to $(|\phi_-^{(1)}\rangle, |\phi_+\rangle, |\phi_-^{(2)}\rangle)$ respectively. Transform to the colour-vertex (flavor) basis via the K3 eigenvector matrix U_{TBM} :

$$\hat{V}_{\text{flavor}}^2 = U_{\text{TBM}} \cdot \text{diag}(16, 144, 900) \cdot U_{\text{TBM}}^T = \begin{pmatrix} 58.\bar{6} & 42.\bar{6} & 42.\bar{6} \\ 42.\bar{6} & 500.\bar{6} & -399.\bar{3} \\ 42.\bar{6} & -399.\bar{3} & 500.\bar{6} \end{pmatrix} \quad (18)$$

(matrix elements are exact rational numbers; decimal display to 1 place after the recurring-digit overline.)

Proposition 5.2 (Exact $\mu\tau$ -exchange symmetry of $\hat{V}_{\text{flavor}}^2$). *Let P_{23} be the permutation matrix swapping rows and columns 2 and 3 (the $\mu \leftrightarrow \tau$ exchange). Then $P_{23}\hat{V}_{\text{flavor}}^2P_{23} = \hat{V}_{\text{flavor}}^2$ exactly. Concretely, $\hat{V}_{22}^2 = \hat{V}_{33}^2 = 500.\bar{6}$, $\hat{V}_{12}^2 = \hat{V}_{13}^2 = 42.\bar{6}$, and the only off-diagonal element distinguishing μ from τ is $\hat{V}_{23}^2 = -399.\bar{3}$, which is symmetric under the swap by being the same on both sides of the diagonal.*

Proof. Direct calculation. The $\mu\tau$ -exchange symmetry of (18) follows from the K3-eigenvector basis structure: $|\phi_+\rangle = (1, 1, 1)^T/\sqrt{3}$ is fully S_3 -symmetric (in particular $\mu\tau$ -symmetric); $|\phi_-^{(1)}\rangle = (2, -1, -1)^T/\sqrt{6}$ is $\mu\tau$ -symmetric (with μ and τ components both equal to $-1/\sqrt{6}$); $|\phi_-^{(2)}\rangle = (0, -1, 1)^T/\sqrt{2}$ is $\mu\tau$ -antisymmetric. The combination $V^2|\phi_+\rangle\langle\phi_+| + V^2|\phi_-^{(1)}\rangle\langle\phi_-^{(1)}|$ contains only $\mu\tau$ -symmetric terms; the $|\phi_-^{(2)}\rangle\langle\phi_-^{(2)}|$ term is $\mu\tau$ -antisymmetric in its eigenvector but $\mu\tau$ -symmetric in the operator (since $(-)\cdot(-) = +$). $\square$

Theorem 5.3 (Exact recovery of TBM angles from $\hat{V}_{\text{flavor}}^2$). *Diagonalizing $\hat{V}_{\text{flavor}}^2$ in (18) returns eigenvalues exactly $(16, 144, 900)$ and eigenvectors exactly $(|\phi_-^{(1)}\rangle, |\phi_+\rangle, |\phi_-^{(2)}\rangle)$ from (2). Recomputing the PMNS mixing angles from these eigenstates returns*

$$\sin^2 \theta_{12} = \frac{1}{3}, \quad \sin^2 \theta_{23} = \frac{1}{2}, \quad \sin^2 \theta_{13} = 0 \quad (19)$$

exactly, matching the SM-5 TBM result (1) at zeroth order.

Proof. Tautological by construction: $\hat{V}_{\text{flavor}}^2$ in (18) is the conjugation of $\text{diag}(16, 144, 900)$ by U_{TBM} ; diagonalizing it inverts that conjugation. Numerical verification by direct computation reproduces (19) to all decimal places [11] §4. $\square$

Theorem 5.3 closes Theorem 5.1 clauses (i) and (ii) at zeroth order. Constraint K1/K2/K3 from the audit [8] §6 is satisfied by construction — the cage-shell mass formula does not displace SM-5's PMNS result; it preserves it exactly under the mass-basis reading.

5.4 The structural-physical question

Theorem 5.1 clause (iii) is the structural-physical question: why does the specific V assignment $(16, 144, 900)$ attach to specific K3 eigenmodes $(|\phi_-^{(1)}\rangle, |\phi_+\rangle, |\phi_-^{(2)}\rangle)$ in that specific order, by 600-cell substrate geometry rather than by fit?

This question is the substantive content of the K3-Cage-Shell Consistency Theorem at the structural level. Without closure on it, the cage-shell assignment is consistent with TBM (Theorem 5.3) but is itself an ansatz inherited from Candidate-C-as-fitted-to-data, not derived from substrate primitives.

Three observations frame the question:

- **The K3 bonding mode is fully symmetric over colour vertices.** $|\phi_+\rangle = (1, 1, 1)^T/\sqrt{3}$ has equal amplitude on $|V_1\rangle, |V_2\rangle, |V_3\rangle$ — the most-symmetric K3 mode under the S_3 permutation symmetry of the K3 graph.
- **The K3 antibonding doublet is degenerate at the K3 level.** $\lambda_- = -1$ has multiplicity 2; any orthonormal pair within this 2D subspace is K3-equivalent. The TBM choice picks specific directions ($\mu\tau$ -symmetric $|\phi_-^{(1)}\rangle$ and $\mu\tau$ -antisymmetric $|\phi_-^{(2)}\rangle$); this selection is what SM-5 ansatzes (Remark 2.1).
- **The cage-shell V values reflect 600-cell shell structure.** $V = 12$ is the icosahedron, the bonded shell of nearest neighbors at the 600-cell vertex; $V = 4$ is the tetrahedral subset; $V = 30$ is the icosidodecahedral shell at $d^2 = 2$. These are 600-cell substrate structures, not K3-internal structures.

The K3-Cage-Shell Consistency Theorem clause (iii), in structural form, is the statement that the embedding of K3 into the 600-cell respects the symmetries that make the V assignment force the specific order ($V = 4, V = 12, V = 30$) to the specific eigenmodes ($|\phi_-^{(1)}\rangle, |\phi_+\rangle, |\phi_-^{(2)}\rangle$). Three candidate routes for the closure are laid out in [11] §6: Route A (symmetry-shell correspondence), Route B (sub-shell decomposition of icosahedron), Route C (direct distance computation from K3 centroid in 600-cell using SM-3 K3 Spectral Theorem). Route C is the most direct anchoring on existing corpus and uses no new postulates; it is the structural-closure path adopted in §5.5.

5.5 Route C: Direct distance computation from the lepton position

Route C anchors on the SM-3 K3 Spectral Theorem [2] (the K3 embedding into the 600-cell as the colour-cage base of the tetrahedral cage) plus the SM-1 four-cage taxonomy [1] plus the 600-cell distance-shell taxonomy used in SM-7–SM-9 [4, 5, 6].

5.5.1 The 600-cell distance-shell structure (verified)

From any 600-cell reference vertex, direct numerical computation of squared distances to all 119 other vertices gives the shell sequence:

Shell index	d^2	Vertex count	SM-1 cage identification
0	0	1	(the reference vertex itself)
1	$1/\varphi^2 \approx 0.382$	12	icosahedral first shell — charm/tau/Z cage
2	1	20	dodecahedral second shell — bottom/Higgs cage
3'	≈ 1.382	12	(intermediate shell, no SM-1 cage assignment)
3	2	30	icosidodecahedral shell — top quark cage; <i>also ν_3 shell</i>
		(further outer shells; complete sequence in [11] §9.1)	
8	4	1	(antipodal vertex)

Table 5: 600-cell distance shells from any reference vertex. Total vertex count is 120 (verified). Shell-1 vertex count of 12 matches SM-1’s icosahedral cage; shell-2 of 20 matches the dodecahedral cage; the $d^2 = 2$ shell of 30 matches the icosidodecahedral cage. **The 600-cell geometry forces the bonded-shell vertex counts to be 12, 20, 30 at the small-V scale** — this is not a free parameter; it is a topological fact about the 600-cell.

The intermediate 12-vertex shell at $d^2 \approx 1.382$ is a non-cage shell (no SM-1 cage uses it). Its vertices form a rotated icosahedron between shells 1 and 2 in distance ordering; SM-1’s cage taxonomy skips it because no stable cage forms at that scale.

5.5.2 $V = 4$ as the tetrahedral subset of shell 1

SM-1 [1] §3 establishes that the *tetrahedral cage* (4 vertices) is “4 of the 12 nearest neighbours in tetrahedral arrangement.” The icosahedron contains regular tetrahedra as the famous compound-of-5-tetrahedra (5 inscribed regular tetrahedra, each using 4 of the 12 icosahedral vertices). The $V = 4$ cage of Candidate C is one such inscribed regular tetrahedron, distinguished by the lepton-flavor structure that picks out specific colour vertices.

This is the $V = 4$ shell of the SF-4 cage-shell assignment (6): not a separate distance shell of the 600-cell, but the **tetrahedral subset of shell 1** at the lepton-cage scale.

5.5.3 $V = 20$ exclusion by SM-1 particle-type taxonomy

A natural question: why does the cage-shell assignment use shells $\{V = 4, 12, 30\}$ for the three neutrino mass eigenstates and not $V = 20$ (the dodecahedral shell 2)?

Per SM-1 [1] §3 and the cage-particle assignment table, $V = 20$ is the **bottom quark and Higgs boson cage**. It is occupied by other particle types in the CPP framework. The neutrino sector, being lepton-flavored, sees the substrate from the lepton position with the lepton-cage structure: $V = 4$ (tetrahedral cage hosting K3 colour vertices), $V = 12$ (full first shell), $V = 30$ (third shell exterior). The $V = 20$ dodecahedral shell hosts particle types whose cage structure has dodecahedral symmetry, which is not the symmetry the K3 lepton-cage structure couples to.

This is not a fit; it is a particle-type taxonomy claim from SM-1 that propagates into SF-4’s cage-shell assignment. **Skipping $V = 20$ for neutrinos is forced, not chosen.**

5.6 The K3-eigenmode-to-shell coupling pattern

With V values $\{4, 12, 30\}$ confirmed as the available cage shells from the lepton position and $V = 20$ excluded by particle-type taxonomy, the remaining question is which K3 eigenmode couples to which shell. Three structural arguments establish the assignment.

Argument 1: $\nu_2 \leftrightarrow V = 12$ forced by symmetry. The K3 bonding mode $|\phi_+\rangle = (1, 1, 1)^T/\sqrt{3}$ has equal amplitude on all three K3 colour vertices, which are 3 of the 4 vertices of the $V = 4$ tetrahedral cage. The 4 cage vertices in turn are 4 of the 12 vertices of the $V = 12$ icosahedral first shell. The bonding mode’s full S_3 -symmetric support across colour vertices generalizes naturally to full H_3 -symmetric support across the icosahedral first shell — the symmetry-hierarchy

$$S_3 \subset H_3 \tag{20}$$

where S_3 acts on the K3 colour vertex base and H_3 acts on the full icosahedral first shell, is the structural basis. The bonding mode “averages” over the three colour vertices of K3 and inherits the H_3 -icosahedral-symmetric global mode of the $V = 12$ shell.

This forces $\nu_2 \leftrightarrow V = 12$ at the structural level. No alternative cage-shell coupling is consistent with full- S_3 symmetry of the bonding mode.

Argument 2: Antibonding modes split between $V = 4$ and $V = 30$. The K3 antibonding modes $|\phi_-^{(1)}\rangle, |\phi_-^{(2)}\rangle$ have non-trivial sign structure across the colour vertices — they break full S_3 symmetry of the K3 base. They cannot couple to the full H_3 -symmetric $V = 12$ shell

as their primary cage; their sign structure has nodes that don't align with any H_3 -symmetric configuration.

The two available alternatives are $V = 4$ (the tetrahedral subset) and $V = 30$ (the icosidodecahedral shell). Both break full H_3 symmetry to subgroups: $V = 4$ to tetrahedral T_d , $V = 30$ to a different subgroup of H_3 (the icosidodecahedron has 30 vertices in 15 antipodal pairs, with rotational symmetry group containing 60 elements). Both are compatible with antibonding-mode sign structure at the symmetry level.

This forces antibonding modes to be split between $V = 4$ and $V = 30$, but does not yet specify which mode goes to which.

Argument 3: The $V = 4$ vs $V = 30$ split inherits SM-5's open problem. The K3 antibonding modes are doubly degenerate at the K3 level alone — any orthonormal basis of the 2D antibonding subspace is K3-equivalent. SM-5 [3] ansatzes a specific basis (the TBM directions $\phi_-^{(1)} = (2, -1, -1)^T/\sqrt{6}$ and $\phi_-^{(2)} = (0, -1, 1)^T/\sqrt{2}$), which is the $\mu\tau$ -symmetric / $\mu\tau$ -antisymmetric pair. SM-5 explicitly registers this selection as an open problem (the lifting of K3 antibonding doublet degeneracy, requiring substrate-internal structure beyond K3 alone).

The K3-Cage-Shell Consistency Theorem clause (iii) does not resolve SM-5's open problem; it inherits SM-5's selection of TBM directions and shows that the V^2 operator's eigendirections in the antibonding subspace coincide with TBM directions exactly (Proposition 5.2). Once SM-5's selection is given, the $V = 4$ vs $V = 30$ split between ν_1 and ν_3 follows from the wavefunction-spread / symmetry-character argument:

- $\nu_1 = (2, -1, -1)^T/\sqrt{6}$ has heavy support on V_1 (the “electron-vertex-asymmetric” mode) — its wavefunction concentrates near the K3 cage itself, $\mu\tau$ -symmetric. Couples to $V = 4$ (the tetrahedral cage hosting the K3).
- $\nu_3 = (0, -1, 1)^T/\sqrt{2}$ has zero support on V_1 — its wavefunction is “outside” the V_1 vertex with antisymmetric support on the $\mu\tau$ pair. Couples to $V = 30$ (the more distant icosidodecahedral shell with antipodal-pair structure).

The argument is rigorous at the symmetry level once SM-5's TBM-direction selection is given: ν_1 has $\mu\tau$ -symmetric character matching the symmetry of the $V = 4$ tetrahedral subset (the tetrahedral cage's three colour-vertex base has $\mu\tau$ -symmetric structure under SM-5's ansatz); ν_3 has $\mu\tau$ -antisymmetric character matching the $V = 30$ icosidodecahedral shell's antipodal pairing structure.

5.7 Route C closure summary and inheritance posture

Route C closes Theorem 5.1 at the following level:

- **V values $\{4, 12, 30\}$ are forced by 600-cell topology**, not fitted (Table 5).
- **$V = 20$ is excluded for neutrinos** by SM-1's particle-type cage taxonomy (§5.5.3).
- $\nu_2 \leftrightarrow V = 12$ is forced by the bonding mode's full S_3 -symmetric character matching the icosahedral first shell via $S_3 \subset H_3$ (§5.6 Argument 1).
- Antibonding modes split between $V = 4$ and $V = 30$ with $\nu_1 \leftrightarrow V = 4$ and $\nu_3 \leftrightarrow V = 30$ via the wavefunction-spread / $\mu\tau$ -symmetry-character argument (§5.6 Argument 3), **inheriting**

SM-5's existing open problem on the lifting of the K3 antibonding doublet degeneracy.

The K3-Cage-Shell Consistency Theorem clause (iii) **introduces no new fitted parameter beyond the inherited single calibration M_0** , but the specific cage-shell-to-K3-eigenmode coupling assignment is a new structural-coupling claim whose theorem-level proof is registered as OPEN-FP-SF-4-2 and is tied to SM-5's existing open problem on the K3 antibonding-doublet-degeneracy lifting. Where SM-5 ansatzes the specific TBM directions in the K3 antibonding doublet (registered as SM-5's open problem), SF-4 inherits that selection and shows that the V^2 mass-operator directions coincide with TBM directions exactly. The cage-shell V values and their assignment to specific eigenmodes are then forced by 600-cell geometry plus the SM-5-ansatzed TBM directions plus the symmetry-hierarchy and wavefunction-spread arguments of §5.6.

Remark 5.4 (Inheritance, not weakening). *SF-4 does not weaken SM-5's open-problem position; it precisely inherits it. Closure of SM-5's open problem on the antibonding-doublet-degeneracy lifting would simultaneously close clause (iii) of Theorem 5.1 to theorem level. They are tied together. This is the same closure level that SM-5 itself operates at: SM-5's TBM = K3-eigenvector result is rigorous given the K3-eigenmode identification ansatz; SM-5 does not derive that ansatz from CPP interaction rules. SF-4's K3-Cage-Shell Consistency, similarly, is rigorous given SM-5's TBM-direction selection in the antibonding subspace; SF-4 does not derive that selection.*

5.8 Open theorem-level work (OPEN-FP-SF-4-2)

The structural-physical picture established in §§5.2–5.7 is at PARTIAL CLOSURE: numerical zeroth-order consistency exact (Theorems 5.3, 5.2); structural closure of clause (iii) at SM-5-inheritance level via Route C; SF-4 introduces no new fitted parameter, but the cage-shell-to-K3-eigenmode coupling assignment is a new structural-coupling claim. Theorem-level closure of OPEN-FP-SF-4-2 at the vertex-by-vertex K3-coupling level is registered as v1.0+ work, with the following sub-goals:

1. Vertex-by-vertex K3-coupling theorem: rigorously establish that the cage-shell coupling pattern is forced not just at the symmetry level (Argument 1, Argument 2) but at the explicit vertex-coupling level for each K3 eigenmode and each cage shell. This requires explicit 600-cell geometric calculation of the K3-eigenmode wavefunction overlap with each candidate shell.
2. Antibonding-doublet split rigorous derivation: closure of the $V = 4$ vs $V = 30$ split for ν_1 vs ν_3 requires resolving SM-5's existing open problem on lifting the K3 antibonding-doublet degeneracy. Tied to SM-5 closure rather than independently achievable in SF-4 scope.
3. Routes A and B as alternative closures: if Route C surfaces an obstruction at theorem-level rigor, fall back to Route A (symmetry-shell correspondence beyond $S_3 \subset H_3$) or Route B (inscribed-tetrahedron sub-shell decomposition with the $V = 4$ tetrahedron as a specific sub-structure of the $V = 12$ icosahedron).

The numerical zeroth-order consistency (Theorem 5.3, exact by construction under mass-basis reading) is robust to which structural route closes Argument 3 to theorem level. The SF-4 paper ships at PARTIAL-CLOSURE level for OPEN-FP-SF-4-2 with explicit pointers to this open work and explicit acknowledgment that theorem-level closure is tied to SM-5's open problem.

6 Predictions Summary

The cage-shell mass formula (5) with the cage-shell assignment (6), the suppression factor (10), and the K3-Cage-Shell Consistency at zeroth order (Theorem 5.3) jointly produce zero-parameter predictions for seven of the eight observable parameters of the active-flavor neutrino sector. The eighth (δ_{CP}) is registered as open and deferred to SF-2 per route (ii) (§7). This section consolidates the predictions for reference.

6.1 Master predictions table

Parameter	SF-4 prediction	Empirical value
<i>Mass eigenvalues (absolute scale: PARTIAL CLOSURE via σ_ν)</i>		
m_{ν_1} (lightest, NH)	0.98 meV	≤ 5 meV (cosmological + Δm^2)
m_{ν_2}	8.81 meV	8.66 meV (from Δm_{21}^2)
m_{ν_3}	55.1 meV	50.9 meV (from Δm_{32}^2)
Σm_ν	64.9 meV	≤ 72 meV (DESI/Planck)
<i>Mass ratios (zero-parameter, V^2 scaling alone — robust to σ_ν closure)</i>		
m_{ν_2}/m_{ν_1}	9.00	8.66
m_{ν_3}/m_{ν_1}	56.25	50.9
<i>PMNS mixing angles (zero parameters at zeroth order; OP-SM-7d for higher order)</i>		
$\sin^2 \theta_{12}$	$1/3 = 0.333$	0.3092 ± 0.0087 (JUNO 2025 first physics; consistent with NuFIT 6.0 [13] value)
$\sin^2 \theta_{23}$	$1/2 = 0.500$	$0.572^{+0.018}_{-0.022}$ (NuFIT 6.0, NO; octant ambiguity)
$\sin^2 \theta_{13}$	0 (0 th order)	0.02203 ± 0.00056 (NuFIT 6.0, NO)
<i>Mass ordering (forced by cage-shell assignment)</i>		
Mass ordering	Normal hierarchy	Unresolved (JUNO multi-year program)
<i>CP phase (deferred)</i>		
δ_{CP}	Open (route ii)	$\sim 195^\circ \pm 100^\circ$

Table 6: Master predictions table for the eight neutrino sector parameters. Seven of eight parameters at zero free parameters; δ_{CP} registered as open and deferred to the EW-sector flagship (SF-2) per route (ii). Empirical values from NuFIT 6.0 [13]; cosmological bound from the DESI/Planck combined analysis as one stringent combination [15, 16], with alternative analyses (e.g., Planck PR4 + supernovae) relaxing the bound to ~ 86 meV [14].

6.2 Structural-residual pattern

The match levels for the predictions distribute into two classes:

- **Tight matches (within 2–4%):** σ_ν at 2%, m_{ν_2} at 2% from Δm_{21}^2 , $m_2/m_1 = 9.00$ at 4% match.
- **Looser matches (8–14%):** m_{ν_3} at 8% from Δm_{32}^2 , $m_3/m_1 = 56.25$ at 11%, $\sin^2 \theta_{12}$ and $\sin^2 \theta_{23}$ at $\sim 10\%$.

The looser matches concentrate on quantities where the SM-5 OP-SM-7d Capotauro corrections (§8) are expected to lift the zeroth-order prediction toward the observed value. The TBM angles deviate from observation by $\sim 10\%$ for θ_{12}, θ_{23} (and $\sin^2 \theta_{13} = 0.02203$ from NuFIT 6.0 is non-zero against TBM's 0); the same Capotauro mechanism that produces those corrections in the EW sector also lifts the V^2 scaling residuals of m_{ν_3} via mode-mode mixing terms beyond leading order.

The tight matches concentrate on quantities not directly affected by Capotauro corrections: σ_ν is purely a substrate-walk-dimension prediction independent of the K3 eigenstructure mixing, and m_{ν_2} corresponds to the K3 bonding mode that is the unmixed singlet (the eigenmode least affected by antibonding-doublet corrections).

6.3 Comparison to discrete-symmetry A_4 models

The mathematical observation that K3 ($\mathbb{Z}_3 \cong A_3$) eigenvectors give the TBM matrix is well-established in the discrete-flavor-symmetry literature [22, 23, 24]. The standard derivation works from A_4 (alternating group on 4 elements) symmetry, identifying TBM as the result of A_4 broken to specific subgroups in charged-lepton and neutrino sectors. The A_4 -from-K3 relation works because $\mathbb{Z}_3 \cong A_3$ is a normal subgroup of A_4 .

The CPP-specific contribution is not the mathematical observation (which is known and not novel) but the *physical identification*: K3 is the colour-cage base of the 600-cell tetrahedral cage in SM-1, not an abstract flavor symmetry postulated ad hoc. SF-4 builds on this physical identification to derive the absolute mass scale ($\sigma_\nu = z^{-10}$ from substrate walk-dimension primitives) and the cage-shell assignment ($V \in \{4, 12, 30\}$ forced by 600-cell topology + SM-1 particle-type taxonomy) — both of which are outside the scope of pure A_4 models, which produce mixing structure but not absolute mass scales without additional Yukawa-coupling assumptions.

The SF-4 framework is therefore complementary to A_4 models: where A_4 explains TBM mixing from group theory, SF-4 grounds TBM mixing in physical substrate geometry and additionally proposes a structurally constrained derivation of the absolute mass scale and the mass-squared splittings from the same substrate primitives at zero free parameters.

7 δ_{CP} Posture (Route ii)

The CP-violating phase δ_{CP} in the PMNS matrix is the eighth observable parameter of the active-flavor neutrino sector. SF-4 does not derive a quantitative prediction for δ_{CP} in v1.0; instead, it adopts route (ii) per [9]: register δ_{CP} as open and defer its derivation to the EW-sector flagship paper SF-2 via OP-SM-7d (Capotauro mechanism) inheritance.

7.1 Justification of route (ii)

The mechanism-selection working document [9] §3 considered two alternative routes for δ_{CP} :

- **Route (i): derive δ_{CP} from CPP primitives within SF-4.** This requires identifying which CPP substrate degree of freedom carries the CP-asymmetric phase information for the neutrino sector and deriving its phase from substrate primitives. The four candidate handles identified in the audit [8] §7 (cage-orientation angle, Capotauro bias in current formalism, K3-eigenstate phase structure, substrate chirality) each suggest distinct multi-session sub-derivation campaigns.
- **Route (ii): defer to SF-2 / EW sector via OP-SM-7d inheritance.** The Capotauro mechanism in the EW sector is the existing CPP machinery that produces higher-order PMNS corrections lifting TBM angles to observed values. The same mechanism that produces those corrections is the natural home for δ_{CP} derivation; its proper paper is SF-2 (electroweak cage bosons), where the EW substrate machinery is the focus.

Route (ii) was selected at Session 39 mechanism review for three reasons:

1. **Scope discipline.** Route (i) introduces a distinct multi-session derivation campaign within SF-4, which is already a heavy-lift paper. Route (ii) keeps SF-4 contained at 7 of 8 zero-parameter predictions plus one explicit deferral.
2. **Inheritance integrity.** OP-SM-7d is already registered in the existing SM-5 corpus [3]; route (ii) preserves that registration and routes it to the appropriate paper (SF-2 EW sector) rather than re-introducing it as a new SF-4 ansatz.
3. **SF-2 readiness.** The eventual SF-2 paper will be the EW flagship and is already scoped to address Capotauro-mechanism corrections to the TBM zeroth-order mixing. Adding δ_{CP} derivation to SF-2's scope at v1.0 drafting is natural; introducing it as a SF-4 sub-derivation would force SF-2 to inherit a SF-4-internal artifact later.

7.2 The four candidate handles for route (i)

For completeness, the four candidate handles for an eventual route-(i) derivation, detailed in [8] §7 and listed here for forward-reference by the SF-2 drafting work:

- **Cage-orientation angle.** The K3 cage embedded in the 600-cell has a discrete set of orientations relative to the substrate's preferred directions; if the cage selects orientations probabilistically with a phase distribution, that phase distribution is a candidate for δ_{CP} .
- **Capotauro bias in the current formalism.** The Capotauro mechanism of OP-SM-7d already introduces a bias correction to TBM angles; if the bias has a phase component, it carries CP information.
- **K3-eigenstate phase structure.** The K3 eigenmodes are real-valued in the basis of (2), but the K3 graph allows complex eigenmodes if one chooses a different basis (e.g., the irreducible representations of $\mathbb{Z}_3$ over the complex numbers, which are the standard cube roots of unity). The phase structure of the complex-irreducible K3 eigenmodes is a candidate handle.
- **Substrate chirality.** The 600-cell substrate has a discrete chiral structure; if neutrino propagation through the substrate carries a chirality-dependent phase, that phase is a candidate.

These remain valid candidates for the eventual SF-2 work. SF-4 v1.0 ships with the explicit deferral; SF-2 v1.0 (subsequent flagship paper) will close the question via whichever handle proves productive.

8 Higher-Order Corrections (OP-SM-7d Inheritance)

The cage-shell mass formula and the K3-Cage-Shell Consistency Theorem deliver zeroth-order results: V^2 scaling for masses (with ~ 4 – 11% structural residuals on the mass ratios) and TBM at zeroth order for the PMNS matrix (with ~ 10 – 14% deviations from the observed angles, plus $\sin^2 \theta_{13} = 0.02203$ from NuFIT 6.0 non-zero against TBM's 0). Higher-order corrections lift both classes of residuals toward observed values. SF-4 inherits these corrections from the existing SM-5 OP-SM-7d (Capotauro mechanism) machinery; it does not derive them.

8.1 Inheritance posture

The Capotauro mechanism is registered in the SM corpus as OP-SM-7d. SM-5 [3] establishes that the higher-order corrections lifting TBM to observed PMNS angles are within the EW-sector substrate machinery (Capotauro bias terms in the K3-eigenstructure expansion). The mechanism’s natural derivation home is the EW-sector flagship (SF-2), which addresses electroweak cage-boson structure including the $W^\pm$, W^0 , Z , and Higgs bosons whose substrate machinery feeds into the Capotauro mechanism.

SF-4’s posture: **register higher-order corrections as conditional on EW-sector closure of OP-SM-7d, do not derive them in SF-4 scope.** This posture is consistent with (a) the SF-line architecture’s separation of concerns (mass-derivation in SF-1, SF-3, SF-4; mixing-corrections via Capotauro mechanism in SF-2), (b) the strict-C inheritance discipline (SF-4 inherits at register-as-open level rather than introducing new ansatz), and (c) the existing SM corpus structure (OP-SM-7d already registered).

8.2 Expected effects of OP-SM-7d closure

The Capotauro mechanism is expected to produce, when closed:

- **TBM angle corrections** lifting the zeroth-order angles
 $\sin^2 \theta_{12}^{(0)} = 1/3, \sin^2 \theta_{23}^{(0)} = 1/2, \sin^2 \theta_{13}^{(0)} = 0$ toward the observed (NuFIT 6.0 [13])
 $\sin^2 \theta_{12} = 0.307, \sin^2 \theta_{23} = 0.572$ (NO, with octant ambiguity), $\sin^2 \theta_{13} = 0.02203$ via mode-mode mixing corrections at $O(\epsilon)$ where ϵ is a substrate-derived small parameter from the EW machinery.
- **Mass-eigenvalue corrections** via the same mode-mode mixing, lifting the structural residuals 4% and 11% in the mass-ratio predictions of (8) toward the observed Δm_{21}^2 and Δm_{32}^2 values.
- **δ_{CP} derivation** as an additional output of the same machinery (per §7 route ii).

The same Capotauro mechanism that lifts TBM to observed angles is therefore expected to also tighten the SF-4 mass-splitting predictions, and to produce δ_{CP} as a derived rather than registered-as-open quantity. Closure of OP-SM-7d in SF-2 would simultaneously close three classes of residuals: (i) TBM mixing-angle deviations, (ii) SF-4 mass-splitting structural residuals, (iii) δ_{CP} open registration.

8.3 What this means for SF-4 v1.0 ship

SF-4 v1.0 ships with explicit acknowledgement that the leading-order V^2 approximation has structural residuals and that the TBM zeroth-order mixing has $\sim 10\%$ deviations. These are not framework failures; they are the precision floor of the leading-order substrate model, with the next-order corrections inherited as conditional on OP-SM-7d closure. The framework’s strength is at the *zero-parameter structural* level; precision at multi-decimal places requires the higher-order corrections that are scope-properly assigned to SF-2.

9 Cumulative Falsifier

The SF-4 framework makes specific zero-parameter quantitative predictions; the predictions are falsified cleanly by experimental measurements that contradict them. This section catalogs five

falsifiers in approximate decreasing order of near-term experimental access. The first three are direct SF-4 falsifiers; the fourth and fifth are framework-level falsifiers that would propagate through SF-4 and other SF-line papers.

9.1 Direct SF-4 falsifiers

9.1.1 Inverted hierarchy (JUNO multi-year program)

The cage-shell assignment (6) forces the mass-squared ordering by the V^2 scaling: $V_{\nu_1} = 4 < V_{\nu_2} = 12 < V_{\nu_3} = 30$ implies $m_{\nu_1} < m_{\nu_2} < m_{\nu_3}$ (normal hierarchy). The framework cannot accommodate inverted hierarchy — no permutation of the cage-shell assignment is consistent with $V \in \{4, 12, 30\}$ from 600-cell topology while reproducing the observed splittings.

The Jiangmen Underground Neutrino Observatory (JUNO) [17] aims to determine the neutrino mass ordering at $\sim 3\sigma$ confidence after several years of data — a multi-year program rather than a single-date deliverable, with sensitivity studies suggesting roughly six years of full operation are required to reach the 3σ threshold via the reactor antineutrino oscillation pattern measurement. **If JUNO determines inverted hierarchy at high confidence, the SF-4 cage-shell assignment is dead**, and the framework’s central derivation chain fails. This is a **clean, named falsifier** with no escape via parameter adjustment — the cage-shell assignment is forced by 600-cell topology, not chosen.

The current global-fit position (NuFIT 6.0 [13]) already shows mild preference for normal ordering, but with sensitivity dependent on inclusion of the Super-Kamiokande atmospheric χ^2 ; without the Super-K atmospheric data, the two orderings are nearly equally well-fit. JUNO’s reactor-antineutrino sensitivity is independent of the atmospheric input and provides the cleanest resolution path. Independent contributions from DUNE [21] and Hyper-Kamiokande over similar timescales will reinforce the JUNO conclusion.

9.1.2 Principled direct-mass falsifier (longer-timescale)

The cage-shell + suppression-factor prediction is $m_{\nu_1} \approx 0.98$ meV (the lightest mass eigenvalue, in the normal-hierarchy convention adopted here). KATRIN [19] currently bounds the beta-decay effective mass $m_\beta = \sqrt{\sum_i |U_{ei}|^2 m_i^2} < 0.8$ eV; future direct-measurement programmes (KATRIN++, Project 8 [20]) target sub-eV sensitivity over the coming decade, with Project 8’s stated programmatic goal at approximately $m_\beta \sim 40$ meV.

Note that direct-mass measurements probe m_β , not m_{ν_1} directly: the relevant observable for SF-4 is the absolute mass scale, which m_β constrains through the PMNS-weighted sum $m_\beta^2 = \sum_i |U_{ei}|^2 m_i^2$. With normal hierarchy and the SF-4 mass values, the predicted $m_\beta \approx 8.7$ meV (dominated by the $|U_{e2}|^2 m_2^2$ contribution).

The principled falsifier. The SF-4 framework predicts a specific absolute scale for m_β via the cage-shell formula plus the PMNS-weighted sum: $m_\beta \approx 8.7$ meV at zeroth order. **A direct-mass measurement that is robustly inconsistent with this prediction would falsify the absolute-scale prediction.** Concretely, a measurement returning either:

- a robust upper bound substantially below ~ 8.7 meV (e.g., $m_\beta < 3\text{--}5$ meV at high confidence), forcing the SF-4 absolute scale to be too large; or

- a measurement substantially above ~ 8.7 meV (e.g., $m_\beta \gtrsim 30\text{--}50$ meV), forcing the SF-4 absolute scale to be too small,

would falsify the $\sigma_\nu = z^{-10}$ structural prediction. **However, the precision required to discriminate at this level is well below the sensitivity floor of currently-funded direct-measurement programmes:** KATRIN reaches ~ 0.5 eV, Project 8 targets ~ 40 meV. The principled-falsifier precision threshold and the near-term experimental sensitivity differ by roughly a decade.

This falsifier should therefore be read as *conceptually clean but not near-term experimentally accessible* via direct beta-decay measurement. The cosmological-sum bound (§9.1.3) is the more realistic near-term constraint on the absolute mass scale. Future programmes pushing beyond the Project 8 ~ 40 meV target into the few-meV regime would convert this falsifier from principled to programmatic.

9.1.3 Cosmological tightening to $\Sigma m_\nu < 50$ meV

The current DESI/Planck combined bound is $\Sigma m_\nu \leq 72$ meV [15, 16]; SF-4 predicts $\Sigma m_\nu = 64.9$ meV (sum of three predicted masses with normal hierarchy). The prediction sits comfortably under the most stringent combination but with limited margin.

Dataset-dependence of the cosmological bound. The 72 meV figure is one stringent combination of DESI BAO + Planck 2018 likelihoods. Alternative analyses with the Planck PR4 likelihood and supernova datasets relax the bound to ~ 86 meV [14], and the bound depends meaningfully on Planck likelihood-implementation choices, the inclusion or exclusion of specific late-time supernova datasets, and the priors on cosmological parameters. The SF-4 prediction $\Sigma m_\nu = 64.9$ meV is consistent with all current cosmological bounds across the dataset-combination range.

Future tightening as the falsifier. Future cosmological surveys (DESI full-survey results, CMB-S4, LSST) are expected to tighten the bound. **Tightening below ~ 50 meV across robust dataset combinations would falsify the SF-4 prediction directly:** $\Sigma m_\nu = 64.9$ meV cannot be reduced to < 50 meV by higher-order corrections without breaking the splitting structure. The cage-shell assignment forces a specific Σm_ν once the absolute scale is fixed by σ_ν .

A bound tightening to $\Sigma m_\nu < 50$ meV across multiple independent dataset combinations is plausible on a 5–10 year timescale; the falsifier strength scales with how dataset-independent the tightening is.

9.2 Framework-level falsifiers

9.2.1 PMNS deviation from TBM at zeroth order

After SM-5 OP-SM-7d Capotauro corrections are subtracted from observed PMNS angles, what remains should be the TBM zeroth-order pattern $\sin^2 \theta_{12}^{(0)} = 1/3, \sin^2 \theta_{23}^{(0)} = 1/2, \sin^2 \theta_{13}^{(0)} = 0$. If precision PMNS measurement (DUNE [21], Hyper-Kamiokande, JUNO) reveals zeroth-order angles that are not TBM after OP-SM-7d corrections, the K3-eigenmode-identification ansatz inherited from SM-5 is in tension. SF-4’s K3-Cage-Shell Consistency would carry that tension as well, since SF-4 inherits and does not introduce the K3-eigenmode ansatz.

This falsifier is framework-level: it falsifies the chain SM-5 $\rightarrow$ SF-4 (K3 spectral derivation), not just SF-4 specifically. Closure of OP-SM-7d in SF-2 is required before this falsifier can be unambiguously evaluated — an obstruction at the OP-SM-7d closure stage would itself be a framework-level signal, separate from PMNS-direct-measurement falsification.

9.2.2 Substrate-mechanism deviation from $\sigma_\nu = z^{-10}$

The suppression factor $\sigma_\nu = z^{-10}$ has the structural form $z^{-2d_{\text{eff}}}$ with integer $d_{\text{eff}} = 5$ from channel enumeration. If precision data force σ_ν to a value not expressible as $z^{-2d_{\text{eff}}}$ for any small integer d_{eff} (e.g., a value that requires $d_{\text{eff}} = 5.3$ or $d_{\text{eff}} = 4.7$, or a non-power-of- z value), the walk-dimension framework is in tension and the suppression mechanism would need re-derivation.

This falsifier is also framework-level: it would falsify the walk-dimension framework’s applicability to neutrino propagation, which has implications for any other unbound-mode CPP physics that uses the same framework. The z^{-10} form is sharp; a precision direct mass measurement that produces an absolute scale of ~ 0.5 meV (a factor of 2 below the prediction) would be inconsistent with $\sigma_\nu = z^{-10}$ and would force the framework to either accept a different integer d_{eff} (revising channel enumeration) or abandon the walk-dimension framework entirely.

9.2.3 Partial-failure scenarios and modular falsification

The five falsifiers above target structurally distinct claims, and their failure modes decouple in informative ways. The two most illuminating partial-failure scenarios:

- **TBM falsified, mass ratios survive.** If precision PMNS measurement reveals zeroth-order angles inconsistent with TBM (after OP-SM-7d Capotauro corrections), but the mass-ratio predictions $m_2/m_1 = 9$ and $m_3/m_1 = 56.25$ continue to match observation within the V^2 structural-residual band, then the SF-4 cage-shell mass mechanism remains viable while the inherited K3-eigenmode-identification ansatz from SM-5 would require revision. The mass-mechanism and the mixing-mechanism are independently falsifiable; falsification of one does not automatically falsify the other.
- **Mass ratios falsified, hierarchy survives.** If the mass-ratio predictions are forced significantly outside the V^2 structural-residual band (say, m_2/m_1 measured to < 5 or > 15 at high confidence), but the normal hierarchy ordering remains intact, then the cage-shell vertex assignment $\{4, 12, 30\}$ would require revision — perhaps a different shell sequence, but the cage-shell mechanism’s structural basis (mass formula proportional to V^α) remains potentially viable. In this case the cage-shell assignment becomes the failure point rather than the mass-mechanism principle itself.
- **Hierarchy falsified.** JUNO confirmation of inverted hierarchy is the cleanest single-claim falsification: the cage-shell assignment $V_1 < V_2 < V_3$ fails, the cage-shell mass formula’s V^2 scaling produces normal-hierarchy mass ordering by construction, and no internal modification of the framework can accommodate inverted hierarchy without abandoning the cage-shell assignment entirely. This is the cleanest framework falsifier on the experimental near-term horizon.

The decoupling structure makes the framework’s claims modularly falsifiable and provides an informative diagnostic posture: each falsifier targets a structurally distinct claim, and partial failures point to which structural claim has failed and which remain viable. This is by design — the strict-C inheritance discipline ensures that inherited content (TBM, M_0 calibration) carries

inherited risk that is structurally separable from the new content (cage-shell assignment, suppression mechanism) introduced by SF-4.

9.3 What is not predicted

SF-4 does not predict Majorana versus Dirac character (§10). Detection of $0\nu\beta\beta$ decay (the GERDA, CUORE, KamLAND-Zen experiments and successors) would constrain Majorana mass via the effective Majorana parameter $m_{\beta\beta}$, but does not directly falsify or confirm the cage-shell mass derivation. A detection or a tightening null result is informative for the eventual SF-4 v1.0+ work on Majorana-vs-Dirac character but is not a near-term cumulative-falsifier item.

Sterile-neutrino predictions are also outside the active-flavor scope of v1.0 (§10). Sterile-neutrino discovery would require extending the cage-shell assignment beyond $V \in \{4, 12, 30\}$ to additional substrate shells; whether the cage-shell mechanism extends to sterile states or specifically rules them out is a v1.0+ open problem.

10 Open Theorem-Level Work

The SF-4 v1.0 paper ships at PARTIAL CLOSURE level for two registered open problems and identifies a third as a structural-derivation handle. This section consolidates the open theorem-level work for forward-reference and to make explicit what v1.0 does not establish.

10.1 Registered open problems

OPEN-FP-SF-4-1: Picture A formalization

The suppression mechanism’s structural-physical picture is established in §4: three convergent CPP physical pictures (A: two-sided DI-bit exchange; B: two ZBW half-cycles per moment; C: edge-straddling coherent state) all give per-channel suppression $\sigma_{\text{channel}} = 1/z^2$, integer channel enumeration gives $d_{\text{eff}} = 5$, and the combined $\sigma_{\nu} = z^{-10}$ matches empirical at 2%.

Theorem-level closure from CPP axioms A1–A11 is v1.0+ work.

The four enumerated sub-goals (§4.5):

1. Picture A formalization (priority): rigorously establish (a) DI-bit send/receive independence in the unbound regime, (b) channel coherence as the AND of both sides aligning, (c) per-moment $1/z^2$ from these.
2. Independence verification: rule out (or quantify) substrate correlations between send and receive choices.
3. Channel-count rigor: quantify sub-leading corrections to $d_{\text{eff}} = 5$ from partial-binding effects.
4. $\alpha = 2$ closure: theorem-level derivation of the $V^{7/3} \rightarrow V^2$ reduction at the bound/unbound boundary.

The expected closure path uses Picture A as the priority candidate (most CPP-axiomatic, anchors directly on DI-bit exchange as substrate primitive). Pictures B and C remain available as alternatives if Picture A surfaces an obstruction. Estimated effort: 5–10 sessions of focused derivation work plus AI-review iteration.

OPEN-FP-SF-4-2: Vertex-by-vertex K3-coupling theorem

The K3-Cage-Shell Consistency Theorem (Theorem 5.1) closes at SM-5-inheritance level via Route C (§5): clauses (i) and (ii) close exactly at zeroth order; clause (iii) closes structurally with V values $\{4, 12, 30\}$ forced by 600-cell topology, $V = 20$ exclusion forced by SM-1 particle-type taxonomy, and the bonding-mode-to- $V=12$ assignment forced by $S_3 \subset H_3$ symmetry-hierarchy. The antibonding-doublet split ($V = 4$ for ν_1 , $V = 30$ for ν_3) inherits SM-5's existing open problem on lifting the K3 antibonding-doublet degeneracy. **Theorem-level closure of OPEN-FP-SF-4-2 at the vertex-by-vertex K3-coupling level is tied to SM-5's open problem.**

The three sub-goals (§5.8):

1. Vertex-by-vertex K3-coupling theorem: explicit 600-cell geometric calculation of K3-eigenmode wavefunction overlap with each candidate cage shell.
2. Antibonding-doublet split rigorous derivation: tied to SM-5 closure rather than independently achievable in SF-4 scope.
3. Routes A and B as alternative closures if Route C surfaces theorem-level obstruction.

The numerical zeroth-order consistency (Theorem 5.3) is robust to which structural route closes Argument 3 to theorem level. Closure of OPEN-FP-SF-4-2 at full theorem-level rigor requires either: (a) closure of SM-5's antibonding-doublet open problem (which simultaneously closes both SM-5 and OPEN-FP-SF-4-2), or (b) an independent SF-4-internal closure that resolves the $V = 4$ vs $V = 30$ split via different machinery (Routes A or B). The estimated effort is comparable to OPEN-FP-SF-4-1 but with high coupling to SM-5 work.

10.2 Structural-derivation handle: $\alpha = 2$ closure

The $V^{7/3} \rightarrow V^2$ reduction at the bound/unbound boundary (§3.3) is currently a structural argument: the linear-cage-dimension factor $V^{1/3}$ has no operational meaning for unbound modes (no rigid cage), so the bound-mode formula's $V^{7/3} = V^2 \cdot V^{1/3}$ decomposition reduces to V^2 alone. Theorem-level closure of this reduction from CPP axioms A1–A11 is part of OPEN-FP-SF-4-1 (sub-goal 4 above) and is naturally bundled with the suppression-mechanism Picture A formalization, since both close together at the rigorous bound/unbound-boundary derivation.

10.3 Items not addressed in v1.0

- **Majorana versus Dirac character.** The cage-shell mechanism in v1.0 does not specify whether neutrinos are Majorana or Dirac fermions. The mass formula $m_{\nu_i} = M_0 V_{\nu_i}^2 \sigma_\nu$ produces three mass eigenvalues without specifying the fermion-number-violating character; both pictures are consistent with the v1.0 result. Registered as open for v1.0+ work; closure depends on substrate-level CP-conjugation properties of the unbound 3D orbital ZBW configuration.
- **$0\nu\beta\beta$ rate prediction.** Depends on the Majorana question above. If neutrinos are Majorana, the effective Majorana parameter $m_{\beta\beta}$ depends on PMNS structure and Majorana phases (which themselves are unspecified in v1.0). Registered as v1.0+ open conditional on Majorana-vs-Dirac closure.

- **Sterile-neutrino predictions.** The cage-shell assignment uses $V \in \{4, 12, 30\}$ (the available shells from the lepton position with $V = 20$ excluded by SM-1 taxonomy). Whether the cage-shell mechanism extends to additional sterile states in higher 600-cell shells, or specifically rules them out, is a v1.0+ open problem. The 600-cell substrate has additional shells at higher d^2 values (Table 5); some of these could in principle host sterile-neutrino-like configurations, but the SF-4 active-flavor scope does not address this.
- **Radiative corrections and renormalization-group running.** The mass formula (5) is at zeroth-order tree level; SF-4 does not address radiative corrections to the cage-shell mass operator nor RG running of the predicted mass eigenvalues from the substrate scale to the experimental scale. The CPP framework’s natural scale is the substrate scale (no UV cutoff in the conventional sense; the substrate provides the regulator), so RG running concerns the embedding of the CPP-derived mass scale into low-energy effective theory rather than RG flow within CPP. Registered as v1.0+ work, naturally bundled with the SF-2 EW closure of OP-SM-7d which addresses similar mass-mixing-running questions for the charged-lepton and quark sectors.
- **Leptogenesis and baryon-asymmetry origin.** The cosmological baryon-asymmetry-of-the-universe (BAU) and the related question of leptogenesis (whether the matter-antimatter asymmetry originates in lepton-sector CP-violation) are outside the active-flavor scope of SF-4. Resolution of these questions requires the EW-sector CP-violating phase δ_{CP} (deferred to SF-2 per route ii) plus cosmological evolution machinery beyond the SF-line scope. SF-4 makes no claim about leptogenesis viability or BAU magnitude.

10.4 Forward roadmap

The five-item open-work catalog above structures SF-4’s continuation:

- **High priority for v1.0 $\rightarrow$ v1.x:** OPEN-FP-SF-4-1 Picture A formalization. Single-paper continuation work; estimated 5–10 sessions of focused derivation.
- **High priority for SM-5 closure $\rightarrow$ v1.x:** OPEN-FP-SF-4-2 vertex-by-vertex theorem (tied to SM-5 open problem). Cross-sector closure benefits both papers simultaneously.
- **SF-2 closure $\rightarrow$ v2.0:** δ_{CP} derivation via OP-SM-7d (route ii). Adds 8th zero-parameter prediction; ships in SF-2, propagates to SF-4 v2.0 update.
- **Long-term:** Majorana-vs-Dirac, $0\nu\beta\beta$, sterile-neutrino extensions. v2.0+ scope; depends on substrate-level CP-conjugation work plus cosmological/experimental data.

11 Discussion

11.1 The programme-level pattern: structural agreement at integer counts as load-bearing signal

Three flagship-class derivations across the SF-line corpus all show the same pattern: SS-7 [4] delivers twelve $N = Z$ alpha-chain nuclei binding energies to 1.5% RMS (with the worst case at $\leq 4\%$) at zero free parameters; SM-9 [6] delivers the top quark mass to 0.02% with $z = 12$ as the only counting input; SF-4 (this paper) delivers the neutrino absolute mass scale via $\sigma_\nu = z^{-10}$ at 2% empirical match. The structural commonality:

- Predictions take the form $M_0 \cdot (\text{integer-counting factor})$ with the integer-counting factor uniquely determined by 600-cell topology.
- No fitted parameters beyond the single calibration $M_0 = m_e \cdot z/\varphi$.
- Structural agreement at zero parameters within the leading-order substrate model's intrinsic precision (1%–12% across cases).
- Higher-order corrections registered as conditional on adjacent-sector closure work, not introduced as new ansatz to tighten precision.

The methodological stance, registered in [10] §9: **structural agreement at integer counts and substrate primitives is the load-bearing signal; precision agreement at multi-decimal places is downstream and framework-idealization-limited**. The validation is not "did we fit observed data to within 0.1%?" but "did we predict the observed integer count and substrate-primitive ratio at zero parameters with all residuals attributable to a small set of registered higher-order corrections?"

This stance distinguishes the SF-line from multi-knob phenomenological models that achieve high precision via free-parameter tuning. The programme's epistemic claim is: when precision agreement at 1%–10% emerges from zero-parameter substrate-topology predictions, the underlying structure is real; precision tightening below 1% requires additional substrate-level mechanisms (the Capotauro-class corrections in OP-SM-7d, the partial-binding corrections in OPEN-FP-SF-4-1) that the framework registers as open derivation work rather than as ad hoc fits.

11.2 Cross-sector implications

SF-4's results have implications for adjacent SF-line papers and the broader CPP programme:

SF-2 (electroweak) closure of OP-SM-7d

OP-SM-7d closure in SF-2 produces three classes of corrections that propagate into SF-4 simultaneously: (i) PMNS angle corrections lifting TBM zeroth-order to observed values; (ii) mass-eigenvalue corrections tightening the V^2 structural residuals (4% and 11% in the mass ratios); (iii) δ_{CP} derivation as an output of the same mechanism. SF-2 closure would extend SF-4's prediction count from 7/8 to 8/8 zero-parameter and would reduce the looser-match residuals in Table 6 from 8–14% toward sub-1%.

SM-5 closure of K3 antibonding-doublet open problem

Closure of SM-5's existing open problem on lifting the K3 antibonding-doublet degeneracy would simultaneously close OPEN-FP-SF-4-2 to theorem level via Argument 3 (§5.6). The two open problems are tied: SF-4 inherits SM-5's ansatz, and theorem-level rigor on the $V = 4$ vs $V = 30$ split for ν_1 vs ν_3 requires the SM-5 selection of TBM directions in the antibonding doublet to itself be derived rather than ansatzed. Cross-sector closure benefits both papers.

Walk-dimension framework: cross-sector applications

The walk-dimension framework introduced in §4.1 (definition of walk channel, walk dimension, $\sigma = N^{-d_{\text{eff}}}$ form) is structurally agnostic about specific particle physics; it applies to any unbound-mode CPP propagation. Candidate cross-sector applications:

- **Free-particle propagators in CPP.** Standard QM free-particle propagators have no suppression at the substrate level; the walk-dimension framework predicts that any CPP free-particle mode carries channel-level coherence dilution. The neutrino is the cleanest case (the bound-mode mass formula explicitly yields zero at $\sigma_\nu = 1$); other unbound modes might carry similar but smaller suppression.
- **Light propagation.** Photons in CPP are unbound modes of the polarized eDP volume; the walk-dimension framework may apply with appropriate channel enumeration.
- **Other gauge bosons.** $W^\pm$ and Z bosons in CPP are cage bosons (bound at the icosahedral and dodecahedral shells respectively); the walk-dimension framework predicts $\sigma_\nu = 1$ for them, consistent with their relatively large masses. This is a successful prediction of the framework, since cage bosons do indeed carry "full mass" without suppression.

These cross-sector applications are pending future investigation and are not in SF-4 scope, but the framework's structural agnosticism suggests one mechanism applied to multiple sectors — a classical mark of programme-level structural correctness.

SF-7 grand unification synthesis

SF-4's 7/8 zero-parameter result is one piece of the cumulative SF-line result. The full SF-line architecture ([12]) covers all 17 Standard Model particles plus the W^0 CPP-novel prediction, with ~ 33 quantitative predictions distributed across the seven flagship papers. SF-4's contribution to the cumulative count: the three neutrino masses (3 predictions), three TBM zeroth-order angles (3 predictions), normal hierarchy (1 prediction) — 7 active-flavor predictions, with δ_{CP} deferred to SF-2.

The eventual SF-7 (grand unification synthesis) paper will sum SF-1 through SF-6 contributions and present the full ~ 33 -prediction cumulative result with all inheritance dependencies traced. SF-4's contribution at v1.0 ship: 7 zero-parameter predictions plus the structural-derivation framework (cage-shell mass formula extended to unbound modes; walk-dimension framework; K3-Cage-Shell Consistency Theorem) that will be reused in SF-1, SF-2, SF-3 with appropriate sector-specific modifications.

11.3 Outlook

The SF-4 framework makes specific zero-parameter quantitative predictions at the structural-numerical level; future precision measurements either confirm or kill it cleanly. The near-term experimental programme provides multiple independent tests:

- **2026–2028:** JUNO mass-ordering determination; the cleanest near-term falsifier (§9.1.1). Inverted-hierarchy confirmation would falsify the cage-shell assignment cleanly.
- **2028–2032:** DESI full-survey results, CMB-S4, LSST cosmological surveys; tightening of Σm_ν bound (§9.1.3).
- **2030+:** KATRIN++, Project 8 direct mass measurements at sub-100 meV sensitivity (§9.1.2).
- **2030+:** DUNE, Hyper-K precision PMNS measurements; framework-level test via TBM-after-OP-SM-7d-corrections (§9.2.1).

The forward research directions on the theoretical side (§10 forward roadmap):

- **Theorem-level closure of OPEN-FP-SF-4-1** (Picture A formalization from CPP axioms A1–A11). Single-paper continuation work; the priority closure path for the suppression-mechanism rigor.
- **Contribution to SM-5 antibonding-doublet open problem.** Closure benefits SM-5 and OPEN-FP-SF-4-2 simultaneously; cross-sector mutual closure.
- **SF-2 EW-flagship drafting** via OP-SM-7d closure. Produces δ_{CP} derivation, higher-order PMNS corrections, and SF-4 mass-splitting precision improvements.

The strategic posture: the SF-4 v1.0 ship is a flagship-class structural-derivation result, with 7 of 8 zero-parameter neutrino-sector predictions in hand. The framework is falsifiable on a 2–10 year experimental timescale. Theorem-level rigor and the missing δ_{CP} prediction are continuation work in the SF-line arc; v1.0 closure is the appropriate ship point for the structural-derivation campaign that opened at Session 37 (the SF-4 audit) and reaches a natural rest point at the v1.0 milestone.

References

- [1] T. L. Abshier, *SM-1: Cage Stability and the Four-Cage Taxonomy of the 600-Cell Substrate*, Conscious Point Physics Standard Model Emergence Series, 2026.
- [2] T. L. Abshier, *SM-3: K3 Spectral Theorem and the Koide Formula*, Conscious Point Physics Standard Model Emergence Series, 2026.
- [3] T. L. Abshier, *SM-5: Tribimaximal Neutrino Mixing from K3*, Conscious Point Physics Standard Model Emergence Series, 2026.
- [4] T. L. Abshier, *SM-7 / SS-7: Alpha-Cluster Polytope Binding Formula and Twelve-Nucleus Verification*, Conscious Point Physics Strong Sector Series, 2026.
- [5] T. L. Abshier, *SM-8: Quark Generation Mass Scaling from 600-Cell Distance Shells*, Conscious Point Physics Standard Model Emergence Series, 2026.
- [6] T. L. Abshier, *SM-9: Symmetry Degeneracy Theorem and the $V^{7/3}$ Mass-Formula Decomposition*, Conscious Point Physics Standard Model Emergence Series, 2026.
- [7] T. L. Abshier, *SS-1: 600-Cell Substrate Foundations and Voronoi Geometry*, Conscious Point Physics Strong Sector Series, 2026.
- [8] SF-4 Neutrino Sector Audit, working document, `flagship_papers/neutrinos/sketches/SF-4_neutrino_sector_audit.md`, Session 37 patch 0294, 2026.
- [9] SF-4 Mechanism Selected (Candidate C with $\alpha = 2$), working document, `flagship_papers/neutrinos/sketches/SF-4_mechanism_selected.md`, Session 39 patch 0298, 2026.
- [10] SF-4 Suppression Derivation, working document, `flagship_papers/neutrinos/sketches/SF-4_suppression_derivation.md`, Sessions 40–41 patches 0299–0300, 2026.
- [11] SF-4 K3-Cage-Shell Consistency Theorem, working document, `flagship_papers/neutrinos/sketches/SF-4_k3_cage_shell_consistency.md`, Sessions 42–43 patches 0302–0303, 2026.

- [12] SF-line Architecture Overview, `flagship_papers/README.md`, Session 41 patch 0301, 2026. Establishes the seven-paper SF-line structure (SF-1 charged leptons, SF-2 electroweak cage bosons, SF-3 quarks, SF-4 neutrinos, SF-5 strong sector, SF-6 electromagnetism, SF-7 grand unification) and the 33-quantitative-prediction scope across all 17 SM particles plus the W^0 CPP-novel prediction.
- [13] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, J. P. Pinheiro, T. Schwetz, and A. Zhou, *NuFit-6.0: Updated global analysis of three-flavor neutrino oscillations*, arXiv:2410.05380 (October 2024); web update at <http://www.nu-fit.org>.
- [14] D. Naredo-Tuero, M. Escudero, E. Fernandez-Martinez, X. Marcano, and V. Poulin, *Living at the edge: A critical look at the cosmological neutrino mass bound*, arXiv:2407.07595 (2024); see also *Neutrino mass bounds from DESI 2024 are relaxed by Planck PR4 and might be further relaxed by ACT*, arXiv:2406.14554 (2024). Documents Planck-likelihood-implementation dependence and supernova-dataset variation in the cosmological Σm_ν bound.
- [15] DESI Collaboration, *DESI 2024 VI: Cosmological Constraints from the Measurements of Baryon Acoustic Oscillations*, arXiv:2404.03002 (2024).
- [16] Planck Collaboration, *Planck 2018 results VI: Cosmological parameters*, Astron. Astrophys. **641**, A6 (2020).
- [17] JUNO Collaboration, *Neutrino Physics with JUNO*, J. Phys. G **43**, 030401 (2016); *JUNO Conceptual Design Report*, arXiv:1508.07166.
- [18] JUNO Collaboration, *First measurement of the θ_{12} mixing angle and Δm_{21}^2 mass-squared splitting from JUNO commissioning data*, arXiv:2511.14593 (2025); first-physics-results values $\sin^2 \theta_{12} = 0.3092 \pm 0.0087$ and $\Delta m_{21}^2 = (7.50 \pm 0.12) \times 10^{-5} \text{ eV}^2$.
- [19] KATRIN Collaboration, *Direct neutrino-mass measurement with sub-electronvolt sensitivity*, Nature Phys. **18**, 160–166 (2022).
- [20] Project 8 Collaboration, *Single electron detection and spectroscopy via relativistic cyclotron radiation*, Phys. Rev. Lett. **114**, 162501 (2015).
- [21] DUNE Collaboration, *Deep Underground Neutrino Experiment (DUNE), Far Detector Technical Design Report*, arXiv:2002.03005 (2020).
- [22] P. F. Harrison, D. H. Perkins, and W. G. Scott, *Tri-bimaximal mixing and the neutrino oscillation data*, Phys. Lett. B **530**, 167 (2002).
- [23] E. Ma and G. Rajasekaran, *Softly broken A_4 symmetry for nearly degenerate neutrino masses*, Phys. Rev. D **64**, 113012 (2001).
- [24] G. Altarelli and F. Feruglio, *Discrete flavor symmetries and models of neutrino mixing*, Rev. Mod. Phys. **82**, 2701 (2010).